



Measuring and Discussing Computer System Performance

or

“My computer is faster than your computer”

The bottom line: Performance

Car	Time to Bay Area	Speed	Passengers	Throughput (pmph)
	3.1 hours	160 mph	2	320
	7.7 hours	65 mph	60	3900

- ° Time to do the task
 - *execution time*, response time, latency
- ° Tasks per day, hour, week, sec, ns. ..
 - *throughput*, bandwidth

Measures of “Performance”

- Execution Time
- Frame Rate
- Throughput (operations/time)
 - Transactions/sec, queries/day, etc.
- Responsiveness
- Performance / Cost
- Performance / Power
- Performance / Energy

How to measure Execution Time?

```
% time program
... program results ...
90.7u 12.9s 2:39 65%
%
```

- Wall-clock time?
- user CPU time?
- user + kernel CPU time?
- **Answer:**

Our definition of Performance

$$\text{Performance}_X \sim \frac{1}{\text{Execution Time}_X}, \text{ for program } X$$

- only has meaning in the context of a **program** or **workload**
- Not very intuitive as an absolute measure, but most of the time we're more interested in **relative performance**.

Relative Performance

- can be confusing
 - A runs in 12 seconds
 - B runs in 20 seconds
 - $A/B = .6$, so A is 40% faster, or 1.4X faster, or B is 40% slower
 - $B/A = 1.67$, so A is 67% faster, or 1.67X faster, or B is 67% slower
- needs a precise definition

Relative Performance (Speedup), the Definition

$$\text{Speedup (X/Y)} = \frac{\text{Performance}_X}{\text{Performance}_Y} = \frac{\text{Execution Time}_Y}{\text{Execution Time}_X} = n$$


Example

- Machine A runs program C in 9 seconds, Machine B runs the same program in 6 seconds. What is the speedup we see if we move to Machine B from Machine A?
- Machine B gets a new compiler, and can now run the program in 3 seconds. Speedup from the new compiler?

What is Time?

CPU Execution Time = CPU clock cycles * Clock cycle time

- Every conventional processor has a clock with an associated clock cycle time or clock rate



What is Time?

CPU Execution Time = CPU clock cycles * Clock cycle time

- Every conventional processor has a clock with an associated clock cycle time or clock rate
- Every program runs in an integral number of clock cycles

Cycle Time

GHz = billions of cycles/second, MHz = millions of cycles/second

Y GHz = 1/Y nanoseconds cycle time

How many clock cycles?

Number of CPU clock cycles = Instruction count * Average
Clock Cycles per Instruction (CPI)

Computer A runs program C in 3.6 billion cycles. Program C requires 2 billion dynamic instructions. What is the CPI?

How many clock cycles?

Number of CPU clock cycles = Instruction count * Average
Clock Cycles per Instruction (CPI)

A computer is running a program with $\text{CPI} = 2.0$, and executes 24 million instructions, how long will it run?

All Together Now

CPU Execution Time = CPU clock cycles * Clock cycle time

CPU clock cycles = Instruction count * Average Clock Cycles per Instruction (CPI)

All Together Now

The diagram shows the formula for CPU Execution Time enclosed in a red rectangular box. The formula is:
$$\text{CPU Execution Time} = \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$$
 Arrows point from each term to its units: 'seconds' points to 'CPU Execution Time', 'instructions' points to 'Instruction Count', 'cycles/instruction' points to 'CPI', and 'seconds/cycle' points to 'Clock Cycle Time'.

seconds

CPU Execution Time = Instruction Count \times CPI \times Clock Cycle Time

instructions

cycles/instruction

seconds/cycle

$$\text{CPU Execution Time} = \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$$

- IC = 4 billion, 2 GHz processor, execution time of 3 seconds. What is the CPI for this program?
- Suppose we reduce CPI to 1.2 (through an architectural improvement). What is the new execution time?

An aside...

Cycle Time/Clock Rate

- Increasingly, modern processors can execute at multiple clock rates (cycle times).
- Why?
- However, the granularity at which we can change the cycle time tends to be fairly coarse, so all of these principles and formulas still apply.

Who Affects Performance?

$$\text{CPU Execution Time} = \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$$

- programmer
- compiler
- instruction-set architect
- machine architect
- hardware designer
- materials scientist/physicist/silicon engineer

Performance Variation

$$\text{CPU Execution Time} = \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$$

	Number of instructions	CPI	Clock Cycle Time
Same machine different programs			
same programs, different machines, same ISA			
Same programs, different machines			

Other Performance Metrics

- MIPS
- MFLOPS, GigaFLOPS, PetaFLOPS (million, billion, quadrillion Floating Point Operations/second)

MIPS

MIPS = Millions of Instructions Per Second

= Instruction Count

Execution Time * 10^6

= Clock rate

CPI * 10^6

- program-independent
- deceptive

Which Programs?

- peak throughput measures (simple programs)?
- synthetic benchmarks (whetstone, dhrystone,...)?
- **Real applications**

Which Programs?

- peak throughput measures (simple programs)?
- synthetic benchmarks (whetstone, dhrystone,...)?
- Real applications
- SPEC (best of both worlds, but with problems of their own)
 - *System Performance Evaluation Cooperative*
 - Provides a common set of real applications along with strict guidelines for how to run them.
 - provides a relatively unbiased means to compare machines.
- Now, we also have more specialized benchmarks for big data, for datacenter applications, for machine learning, ...

Amdahl's Law

- The impact of a performance improvement is limited by the fraction of execution time affected by the improvement

$$\text{Execution time after improvement} = \frac{\text{Execution Time Affected}}{\text{Amount of Improvement}} + \text{Execution Time Unaffected}$$

Amdahl's Law


- The impact of a performance improvement is limited by the fraction of execution time affected by the improvement

$$\text{Execution time after improvement} = \frac{\text{Execution Time Affected}}{\text{Amount of Improvement}} + \text{Execution Time Unaffected}$$

- Make the **common case** fast!! (usually)

Amdahl's Law and Massive Parallelism

Speedup



Key Points

- Be careful how you specify performance
- Execution time = instructions * CPI * cycle time
- Use real applications
- Use standards, if possible
- Make the common case fast