
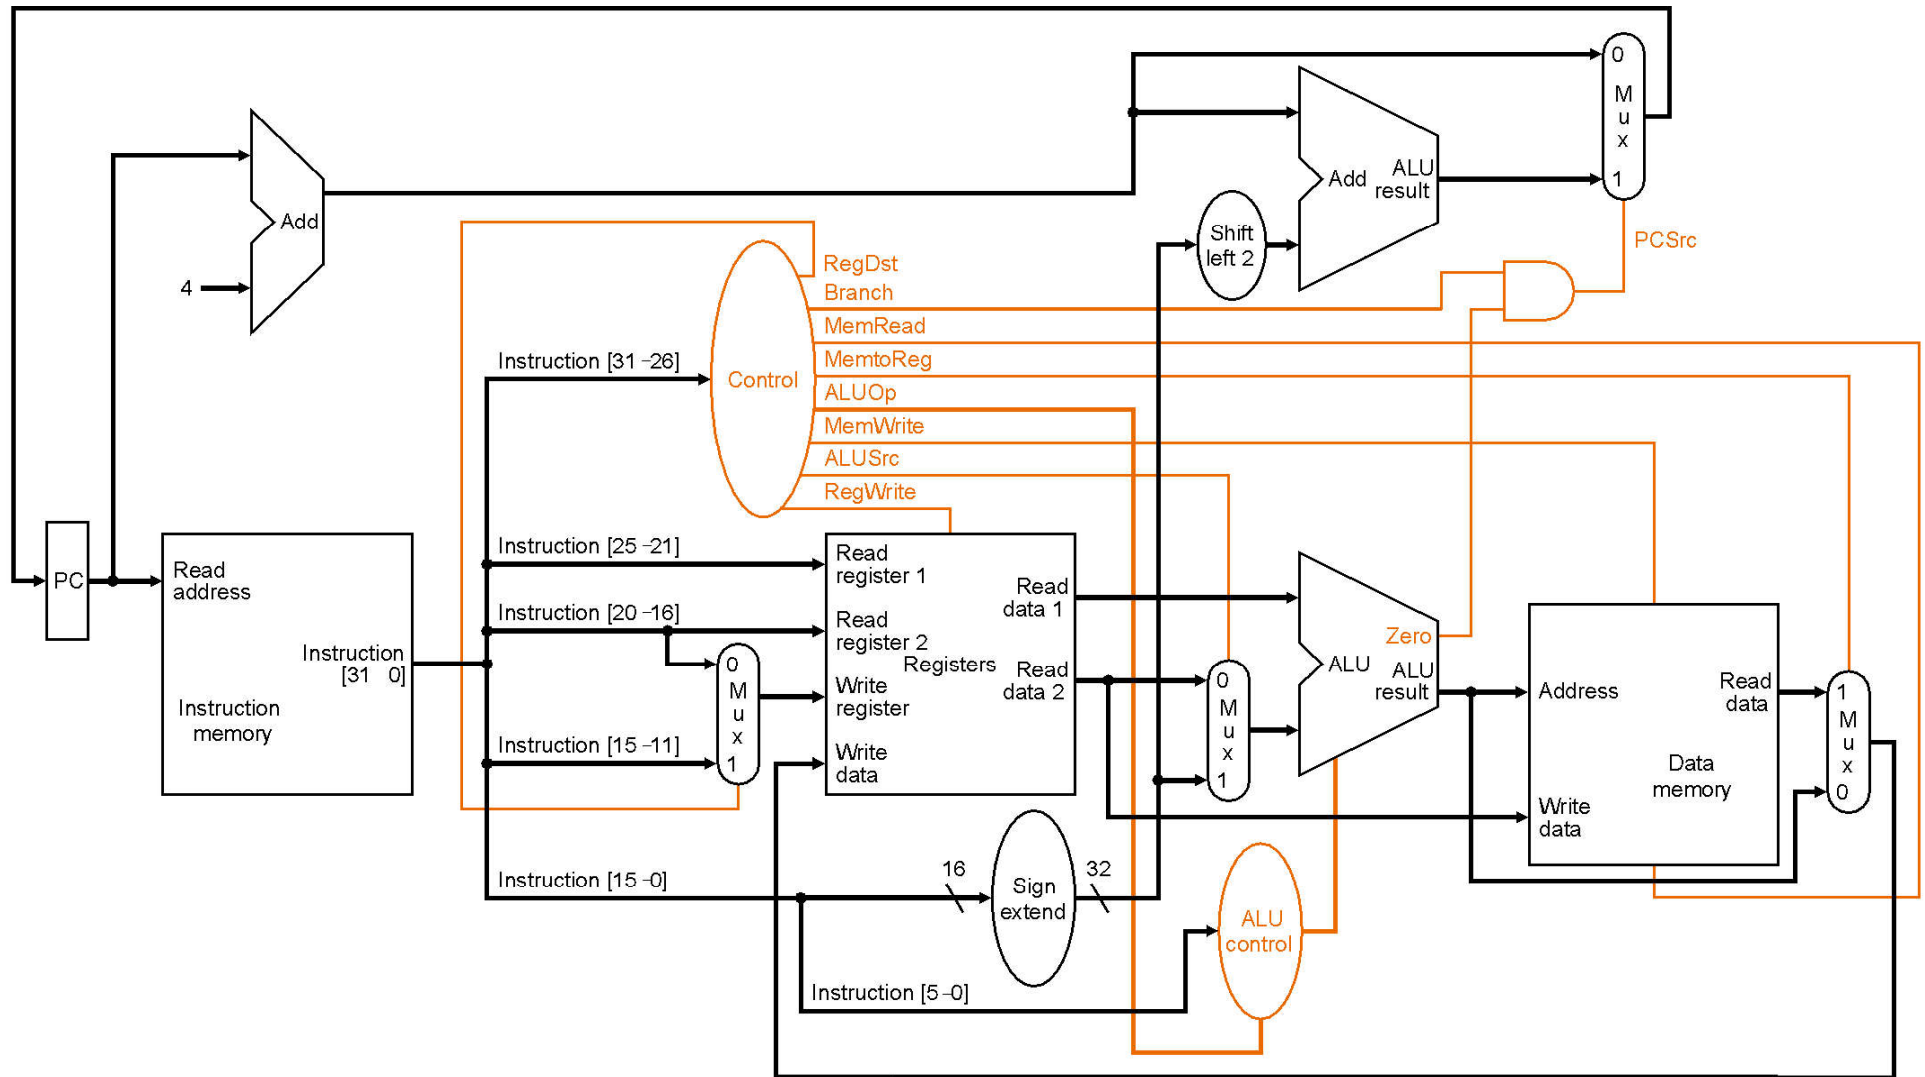


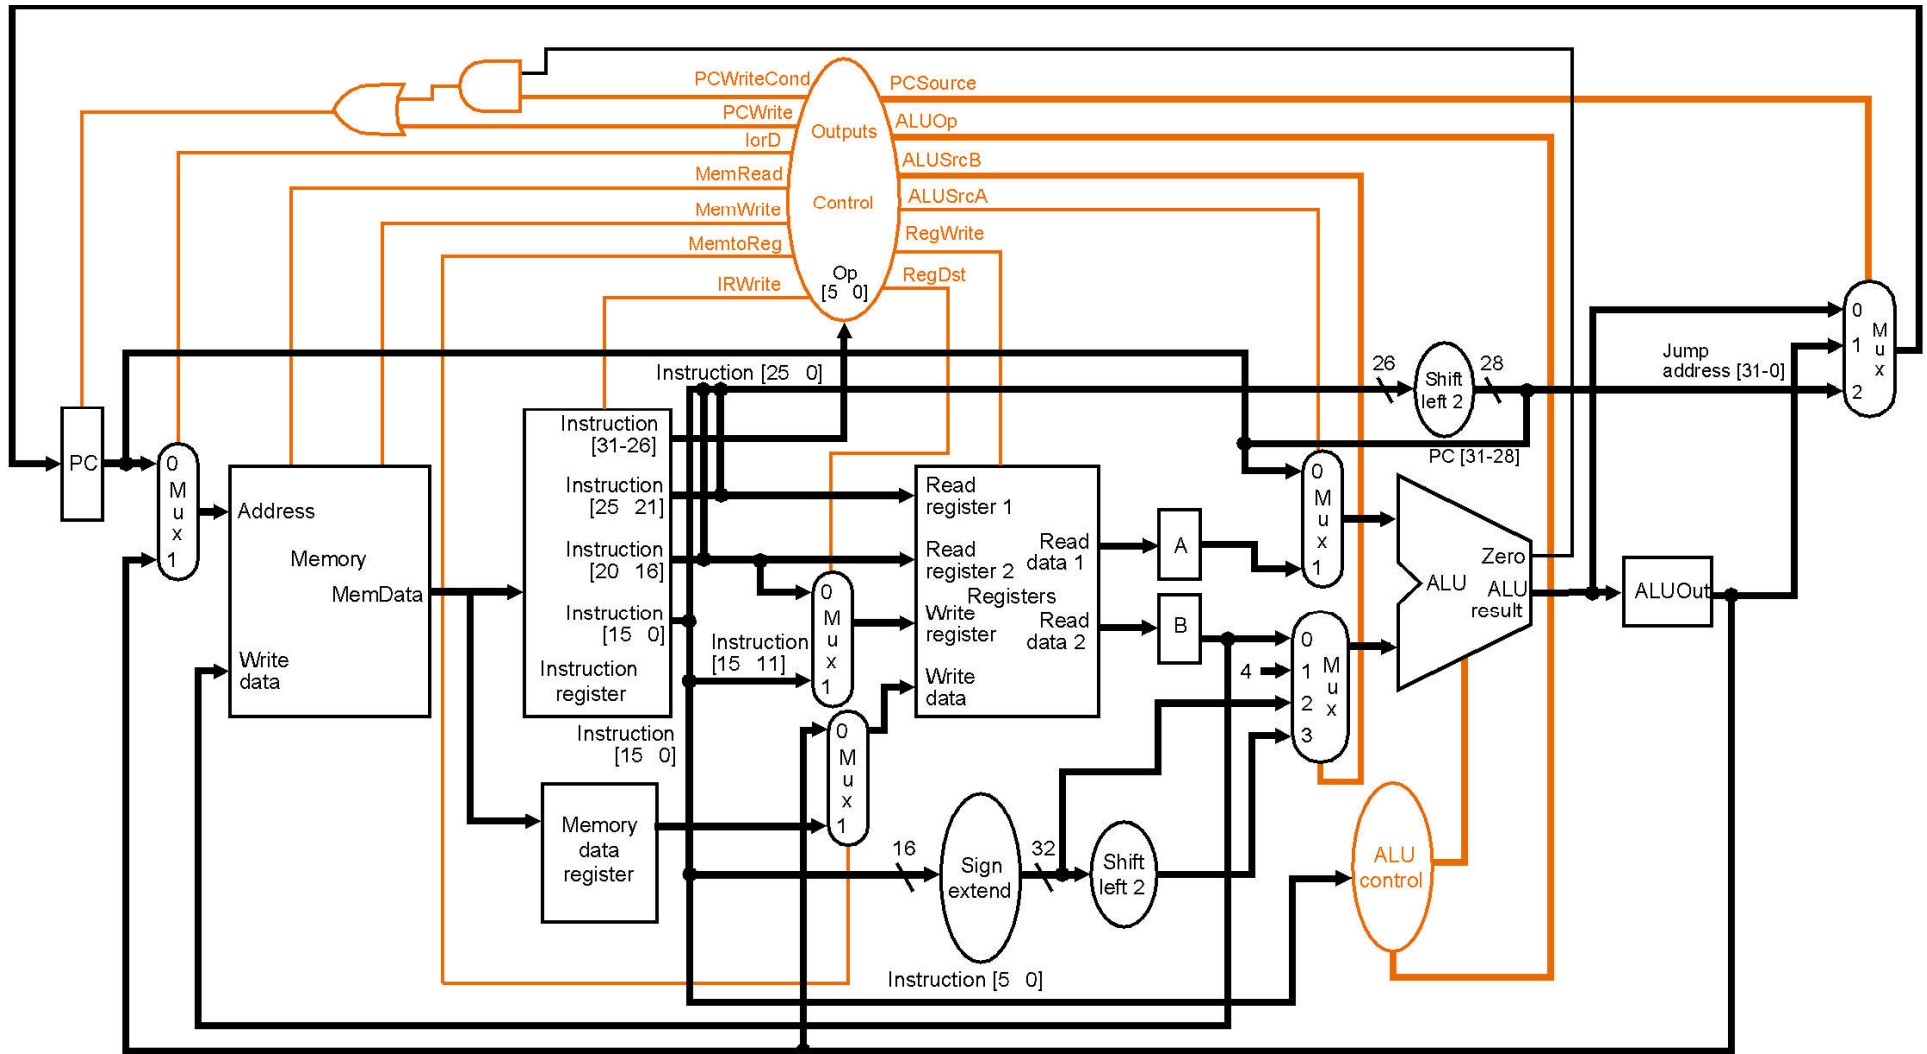
# Designing a Pipelined CPU



# Review -- Single Cycle CPU

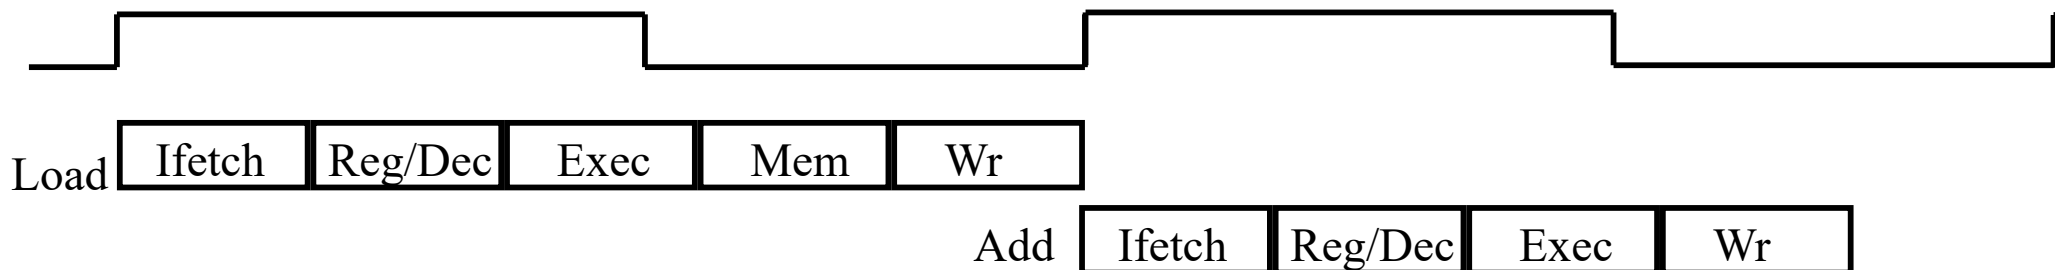


# (not quite) Review -- Multiple Cycle CPU

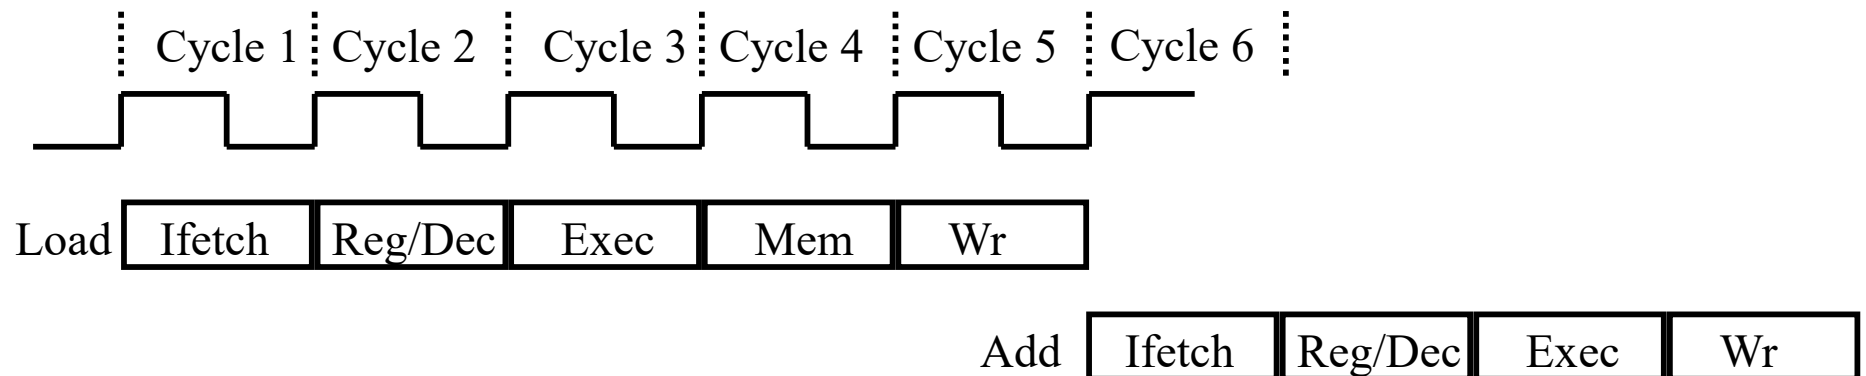


# Review -- Instruction Latencies

## •Single-Cycle CPU




## •Multiple Cycle CPU




# Instruction Latencies and Throughput

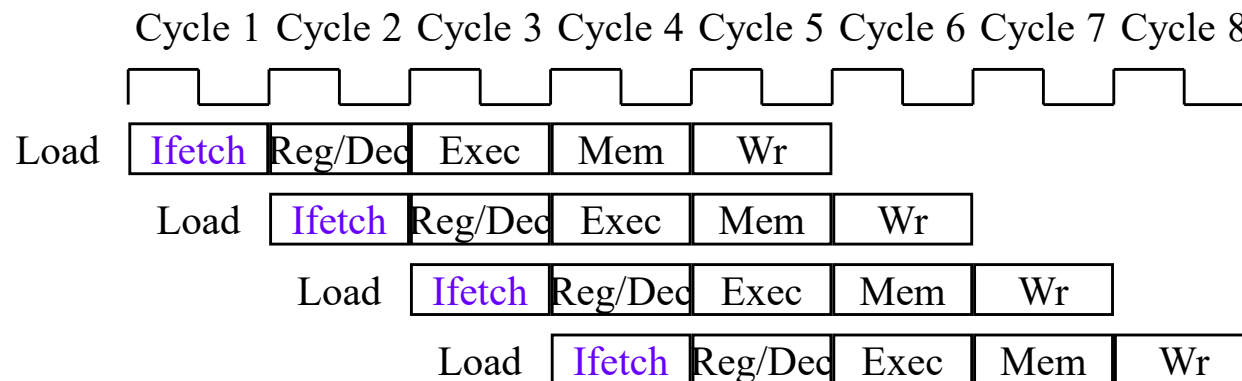
## •Single-Cycle CPU



## •Multiple Cycle CPU



## •Pipelined CPU



# Pipelining Advantages

- Higher *maximum* throughput
- Higher *utilization* of CPU resources
  
- But, more complicated *datapath*, more complex control(?)

# Pipelining Advantages

<u><i>CPU Design Technology</i></u>	<u><i>Control Logic</i></u>	<u><i>Peak Throughput</i></u>
Single-Cycle CPU	Combinational Logic	$\frac{1}{1}$
Multiple-Cycle CPU	Messy (state machine)	$\frac{1}{N}$
Pipelined CPU		$\frac{1}{1}$

# Pipelining in Modern CPUs

- CPU Datapath
- Arithmetic Units
- System Buses
- Software (at multiple levels)
- etc...



# A Pipelined Datapath

IF: Instruction fetch

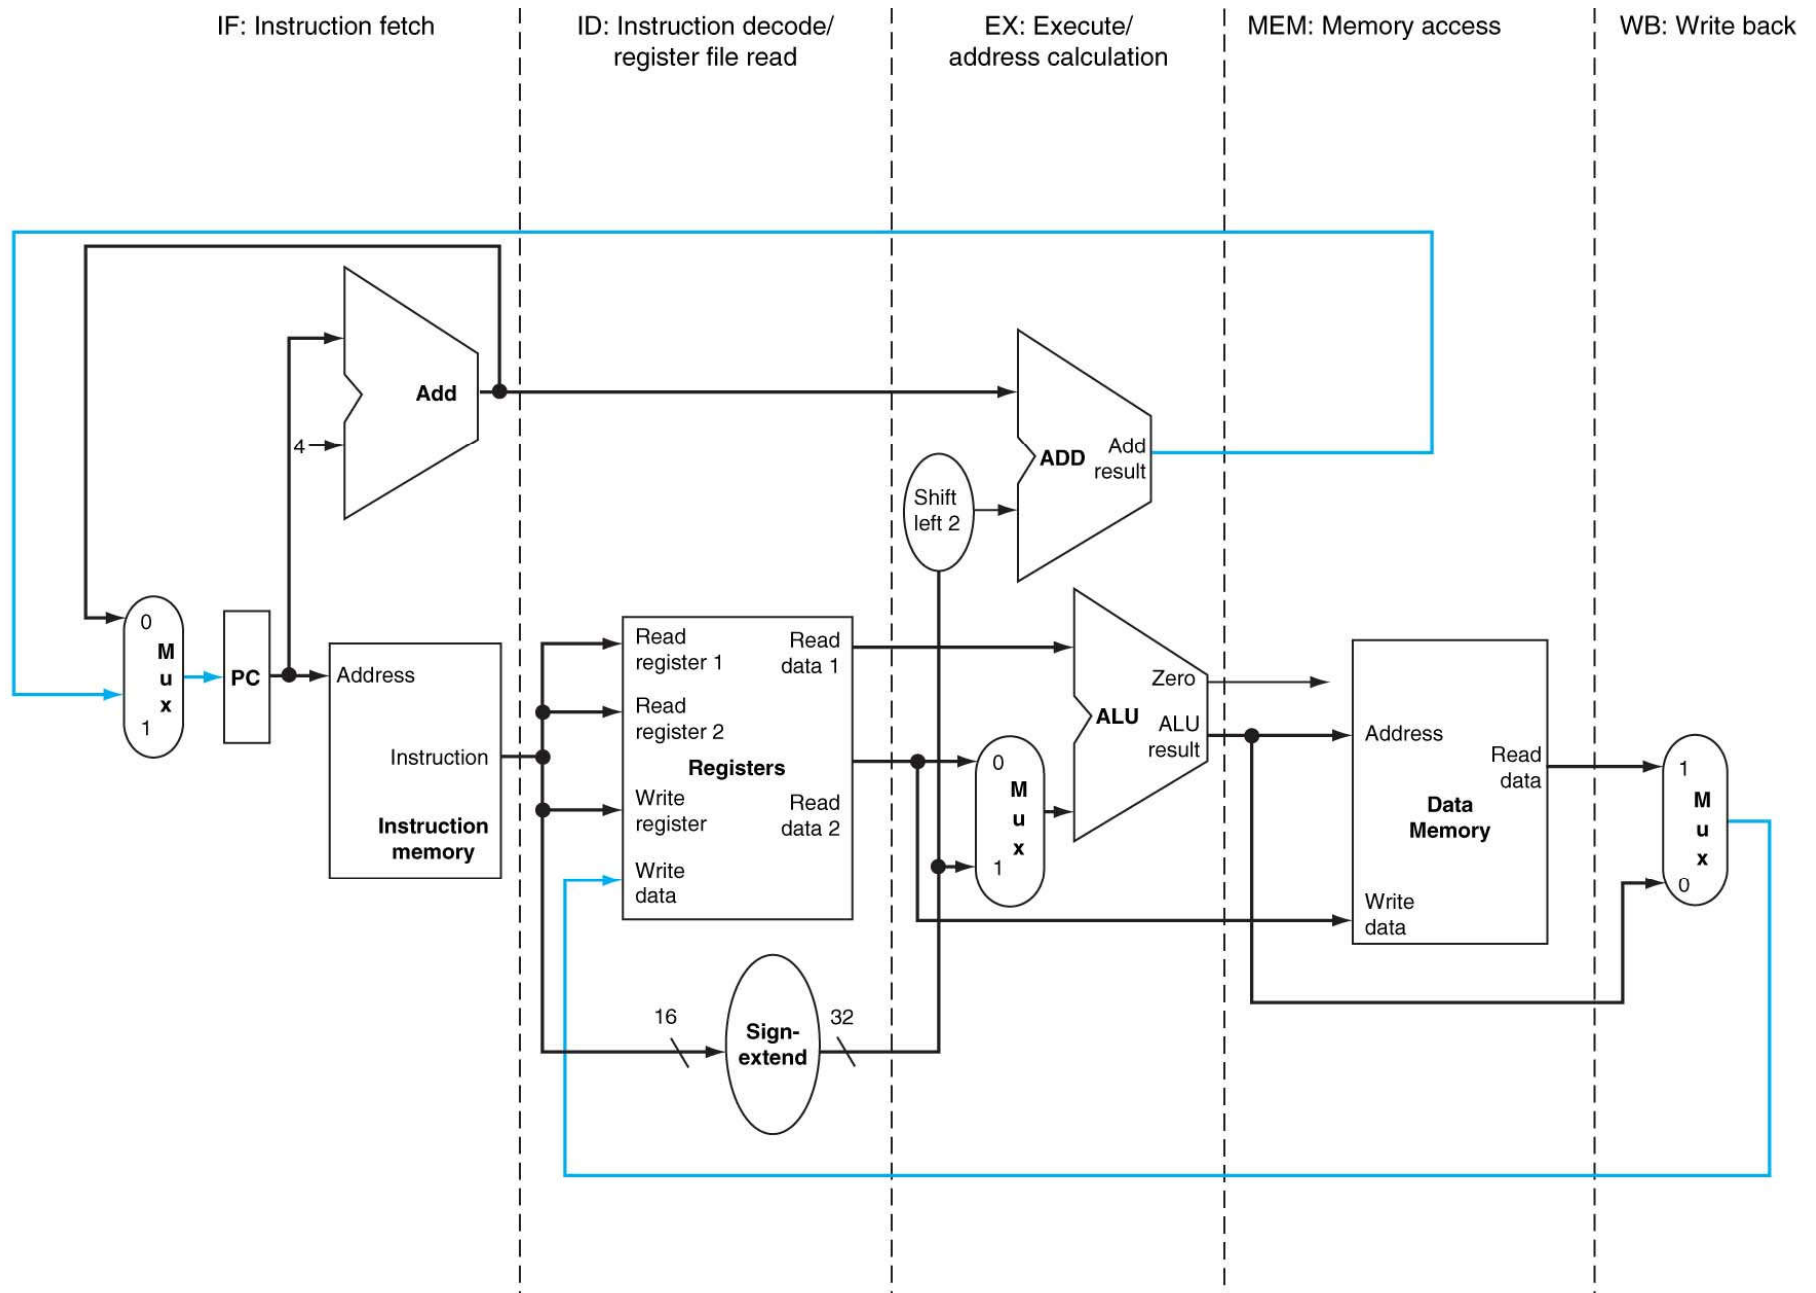
ID: Instruction decode and register fetch

EX: Execution and effective address calculation

MEM: Memory access

WB: Write back

# Pipelined Datapath (roughly)



# Pipelined Datapath

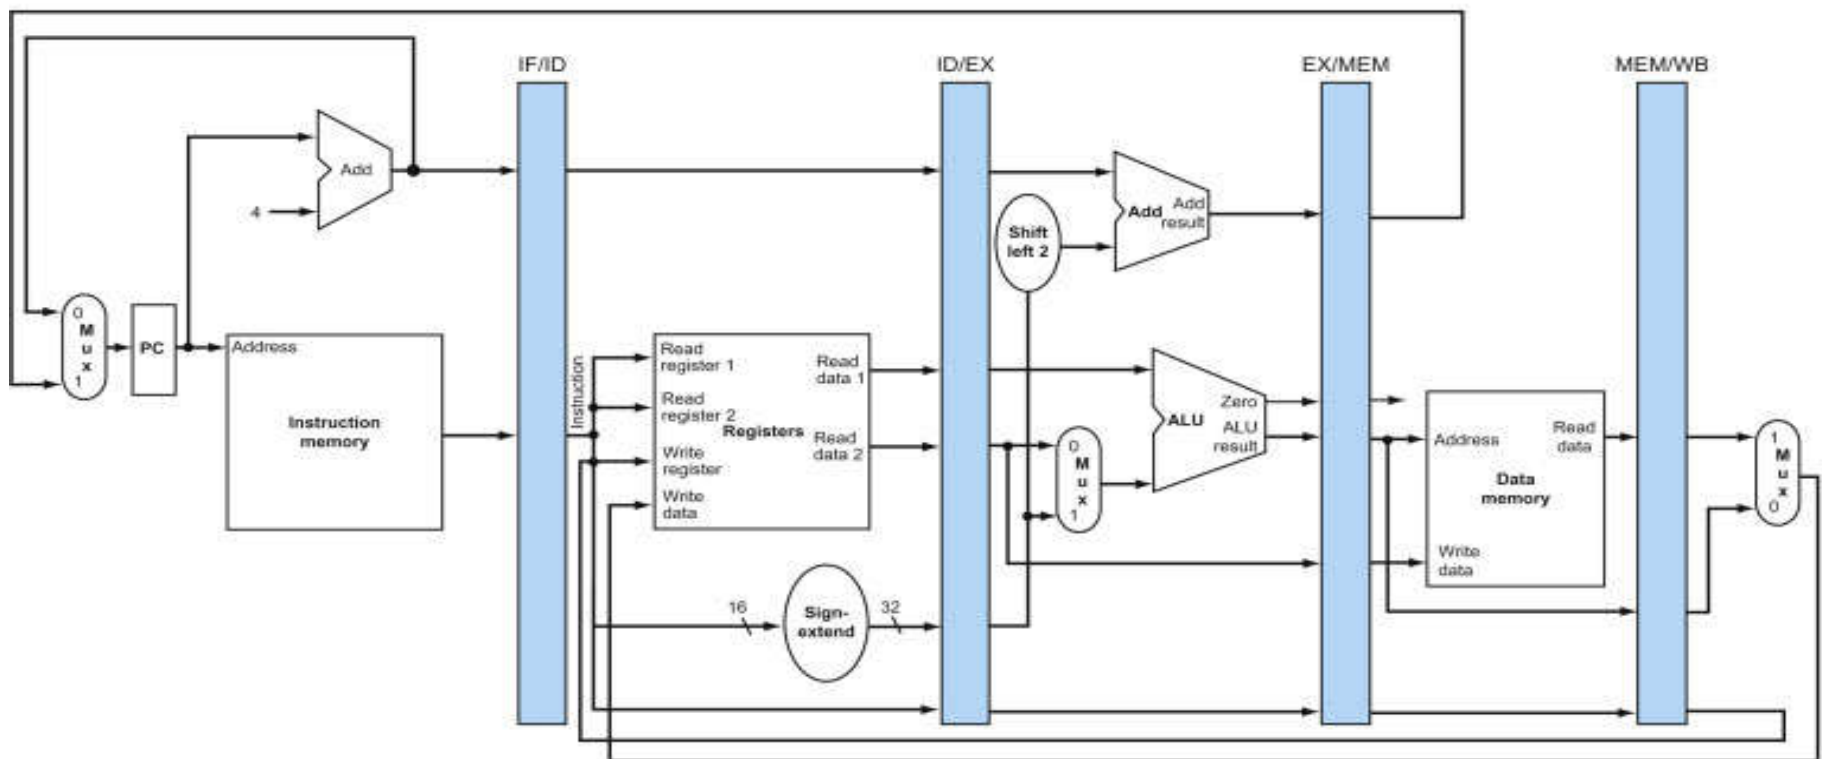
Instruction Fetch

Instruction Decode/  
Register Fetch

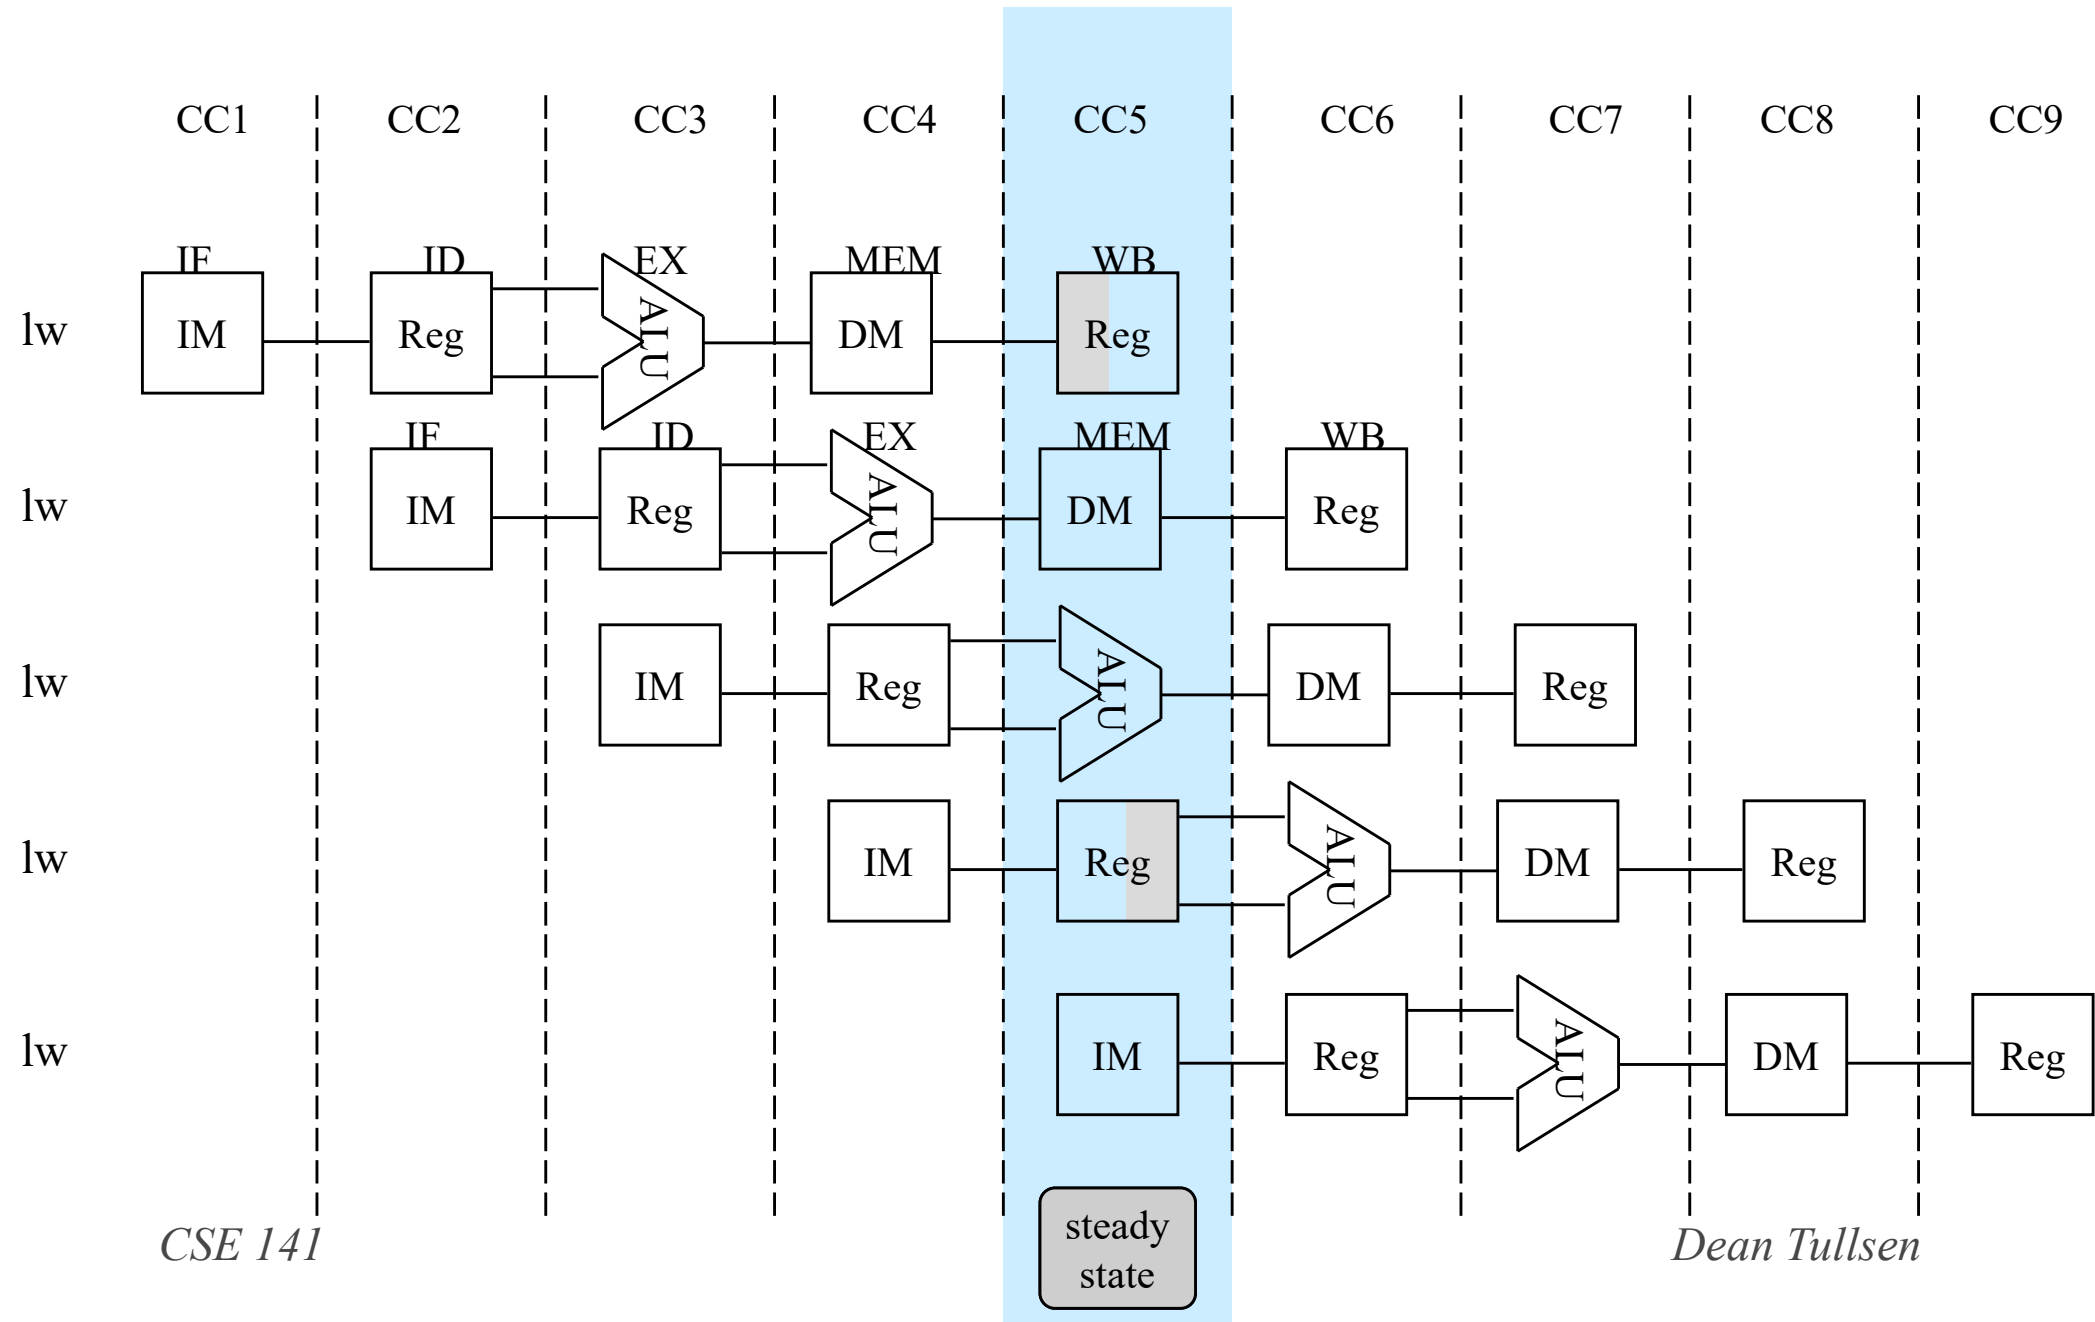
Execute/  
Address Calculation

Memory Access

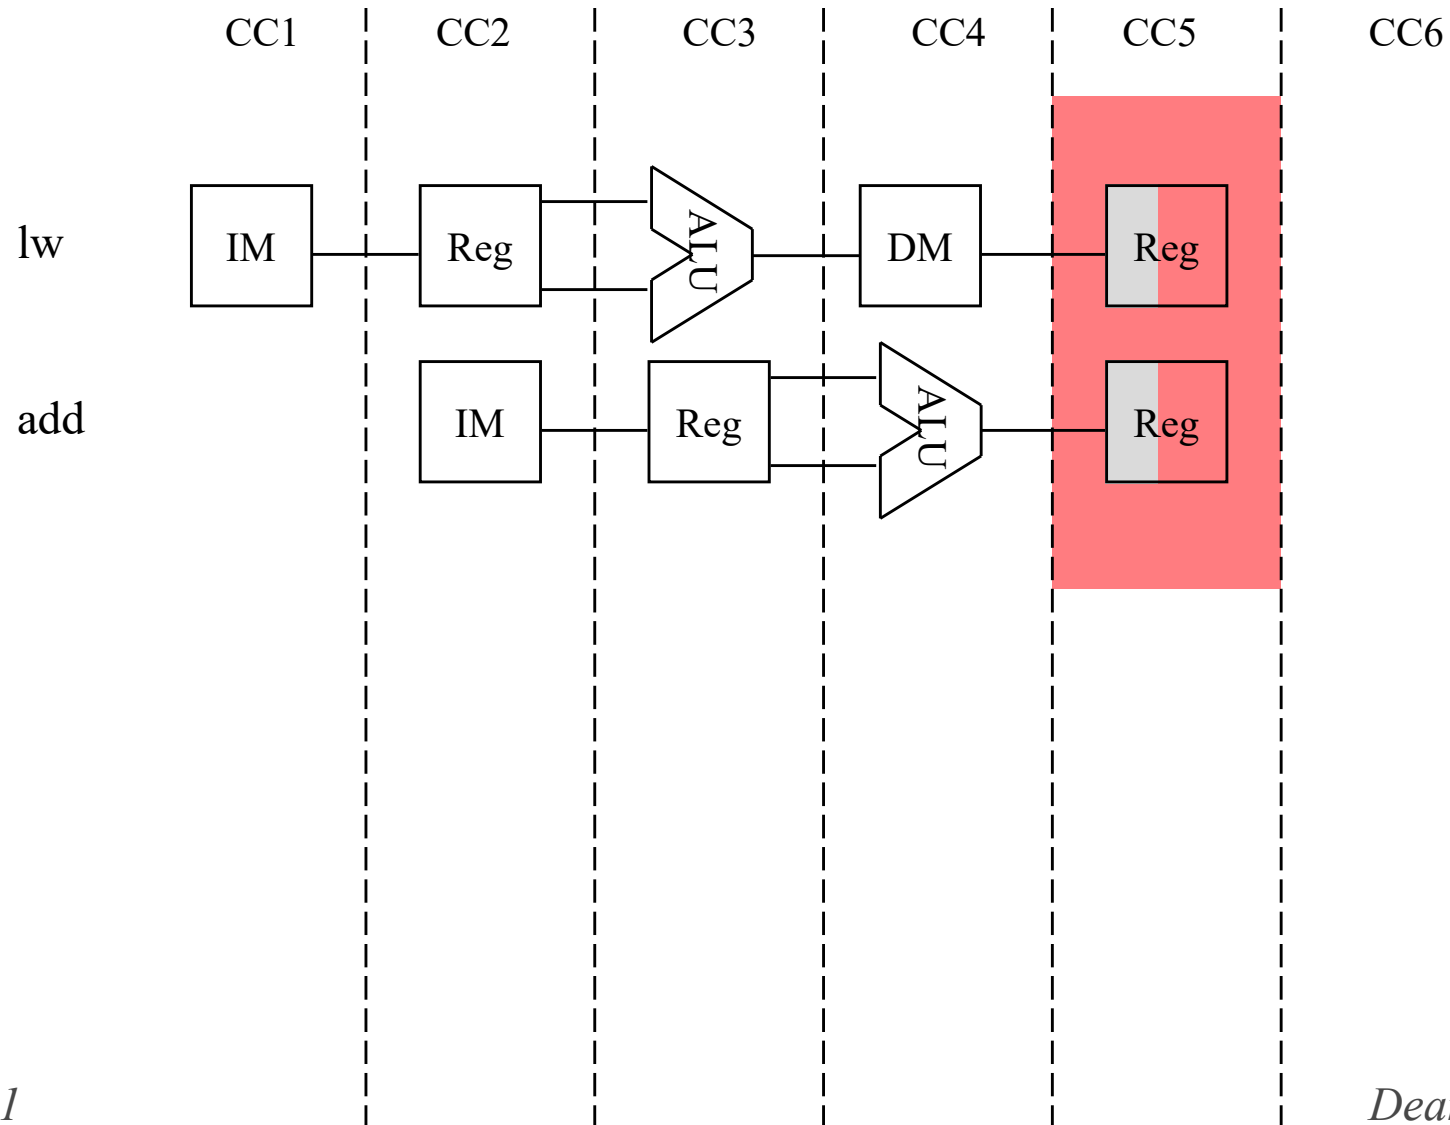
Write Back



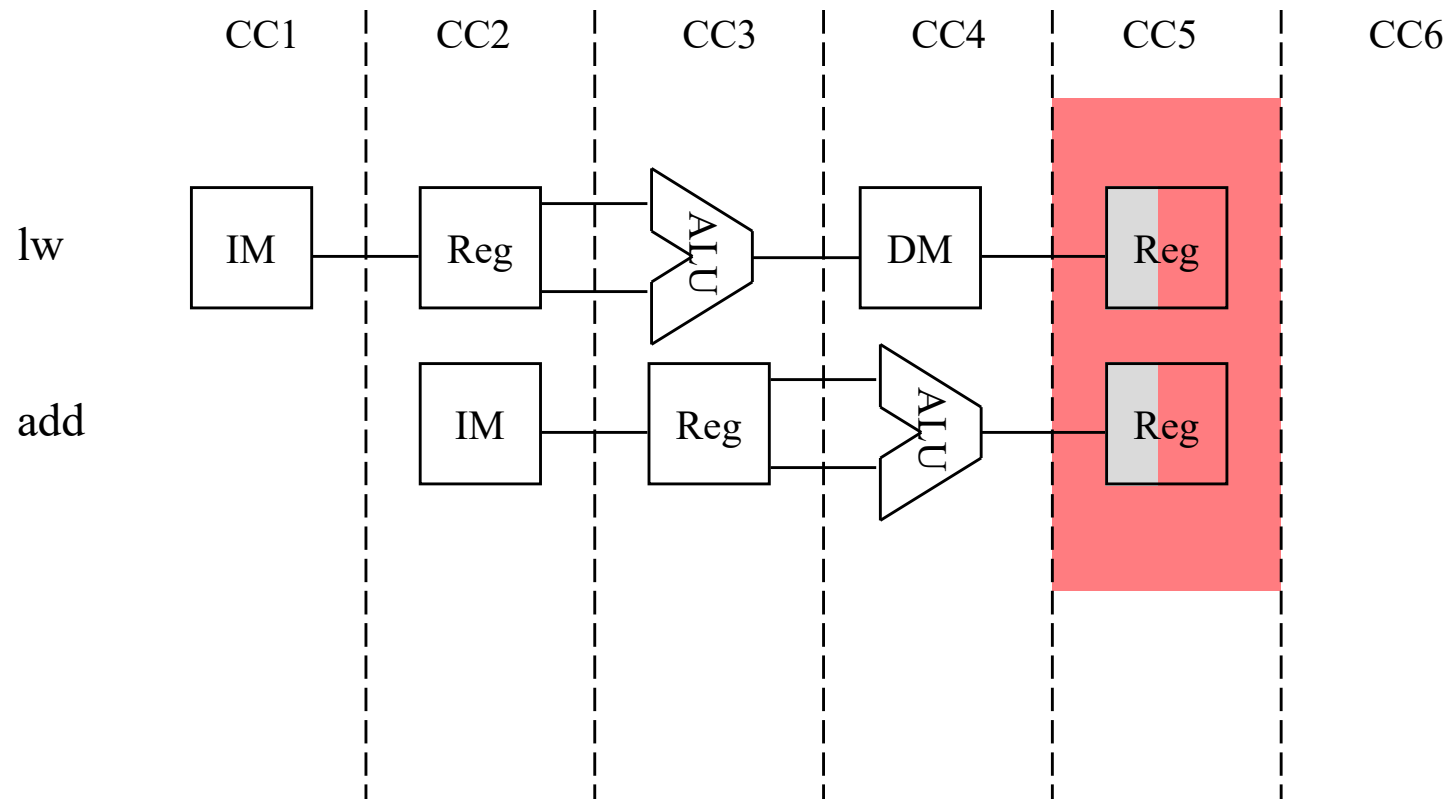
# Execution in a Pipelined Datapath



# Mixed Instructions in the Pipeline



# Mixed Instructions in the Pipeline




This is called a **structural hazard** – too many instructions want to use the same resource.

In our pipeline, we can make this hazard disappear (next slide). In more complex pipelines, structural hazards are again possible.

# Pipeline Principles

- All instructions that share a pipeline should have the same *stages* in the same *order*.
  - therefore, *add* does nothing during Mem stage
  - *sw* does nothing during WB stage
- All intermediate values must be latched each cycle.



# Pipeline Stages

What is the performance implication of making every instruction go through all 5 stages? (e.g., instead of 4 for add, 3 for beq, etc.)

Selection	(Choose BEST answer)
A	Decreases peak throughput by 20%
B	Increases program latency by 20%
C	No significant impact on peak throughput or program latency
D	Depends on how many R-type instructions, beq, etc.
E	None of the above



# Pipelined Datapath

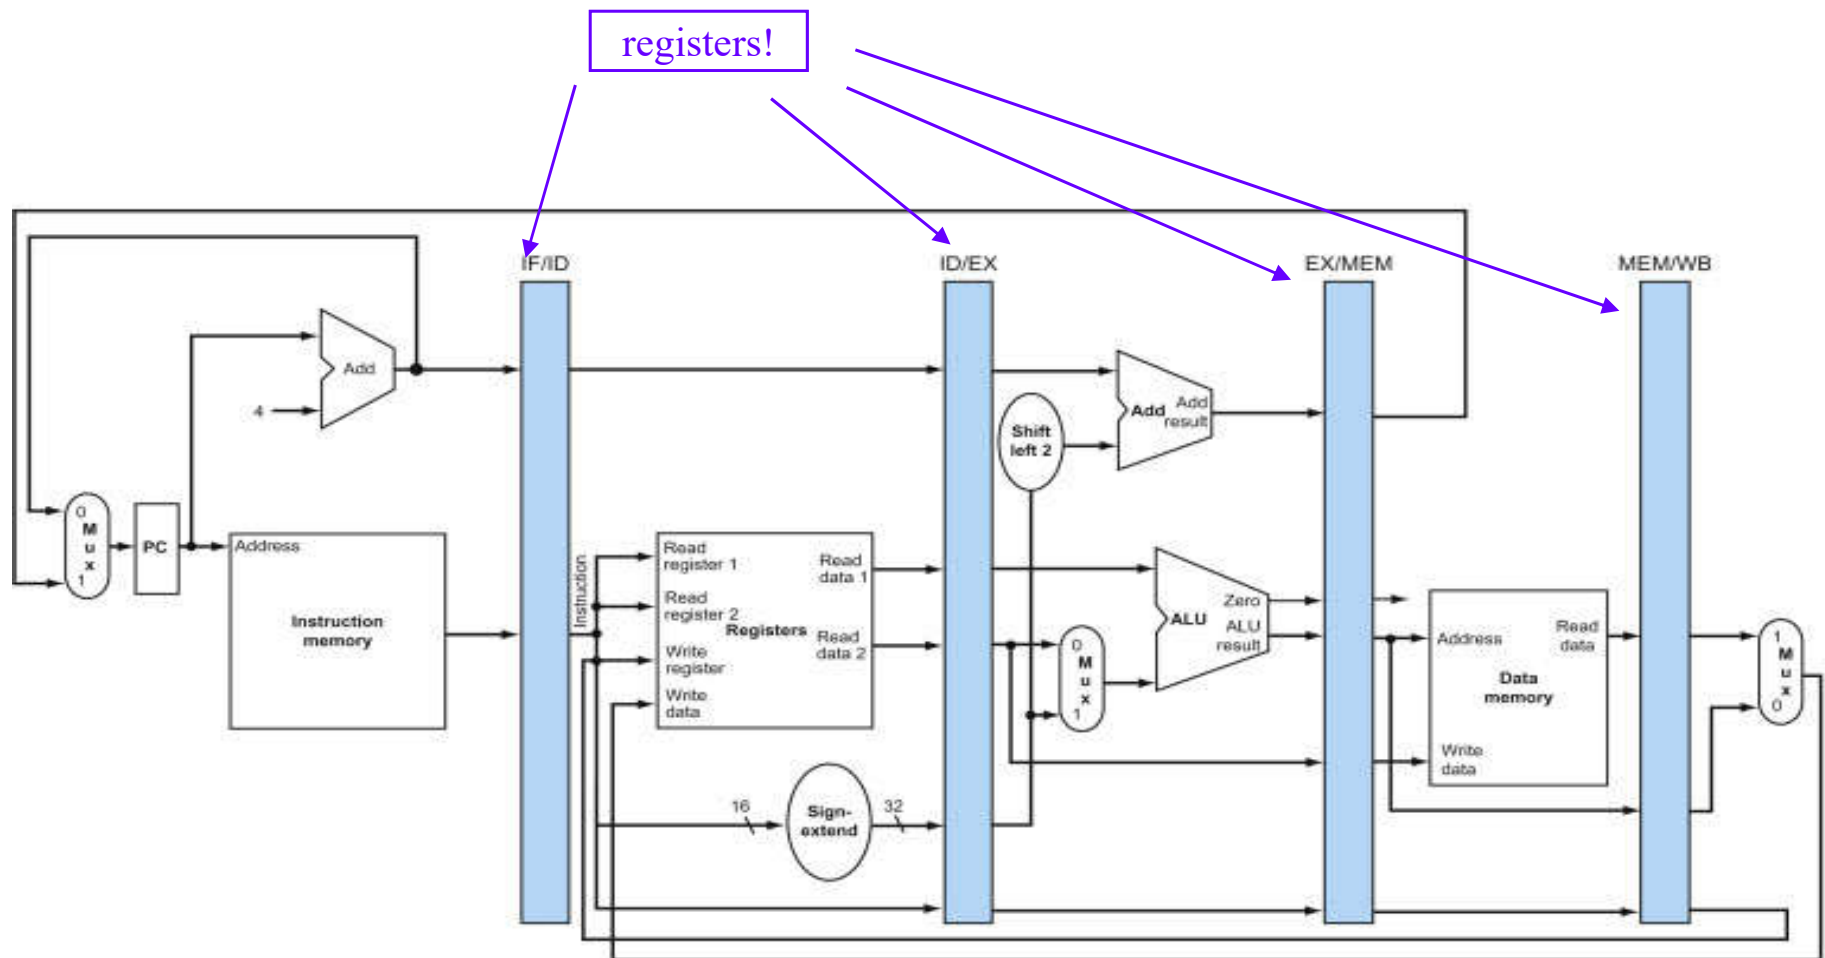
Instruction Fetch

Instruction Decode/  
Register Fetch

Execute/  
Address Calculation

Memory Access

Write Back



# The Pipeline in Execution


**add \$10, \$1, \$2**

Instruction Decode/  
Register Fetch

Execute/  
Address Calculation

Memory Access

Write Back



# The Pipeline in Execution


**lw \$12, 1000(\$4)**

**add \$10, \$1, \$2**

Execute/  
Address Calculation

Memory Access

Write Back



# The Pipeline in Execution


sub \$15, \$4, \$1

lw \$12, 1000(\$4)

add \$10, \$1, \$2

Memory Access

Write Back



# The Pipeline in Execution


Instruction Fetch

**sub \$15, \$4, \$1**

**lw \$12, 1000(\$4)**

**add \$10, \$1, \$2**

Write Back



# The Pipeline in Execution

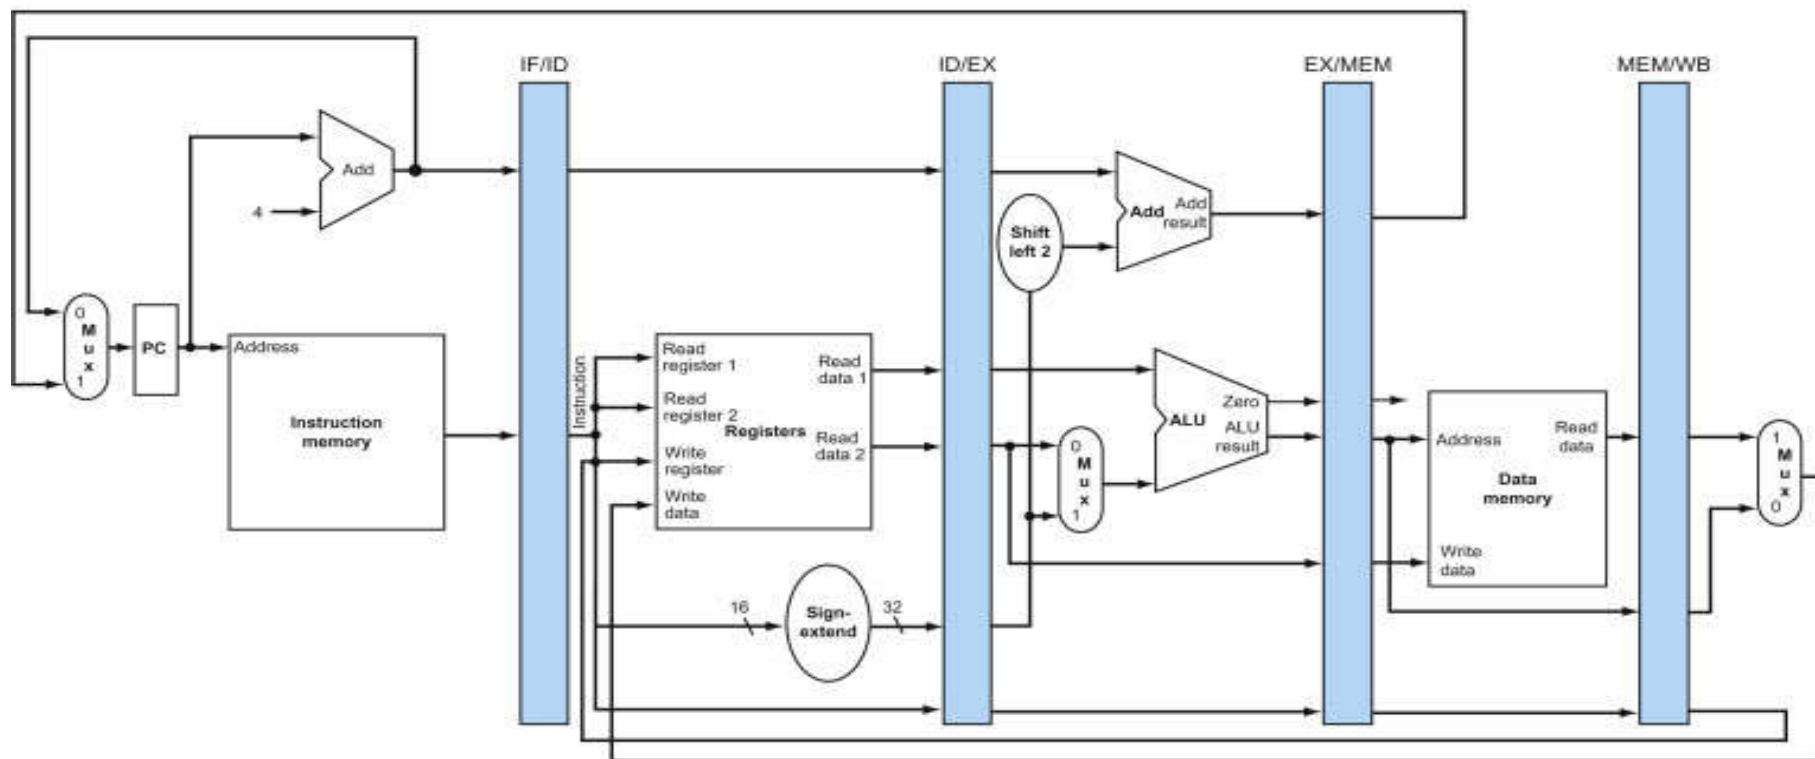
Instruction Fetch

Instruction Decode/  
Register Fetch

**sub \$15, \$4, \$1**

**lw \$12, 1000(\$4)**

**add \$10, \$1, \$2**



# The Pipeline in Execution

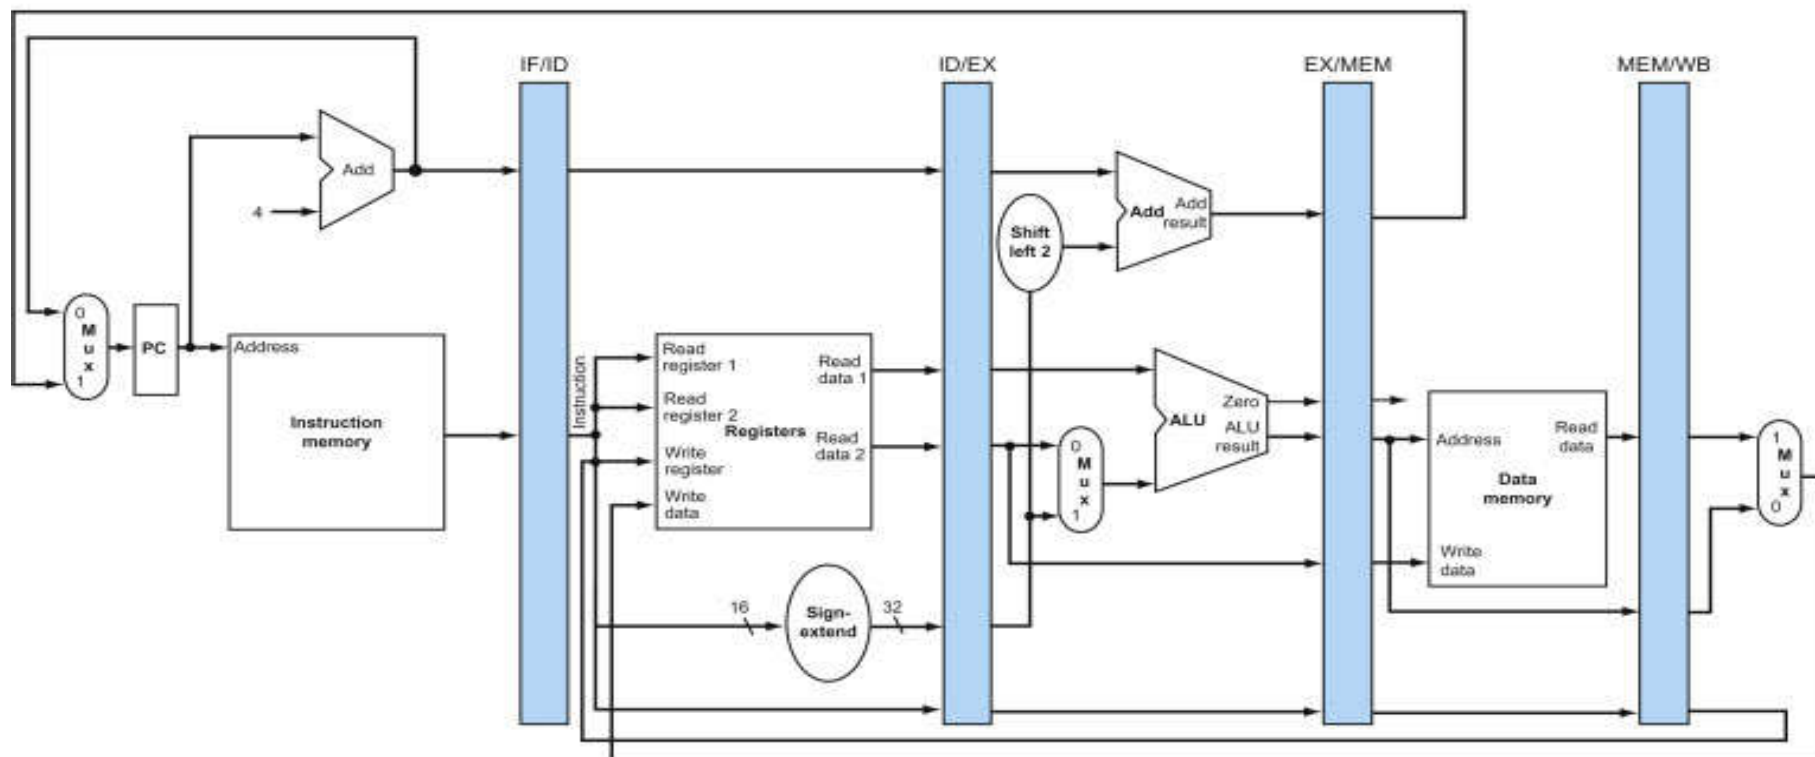
Instruction Fetch

Instruction Decode/  
Register Fetch

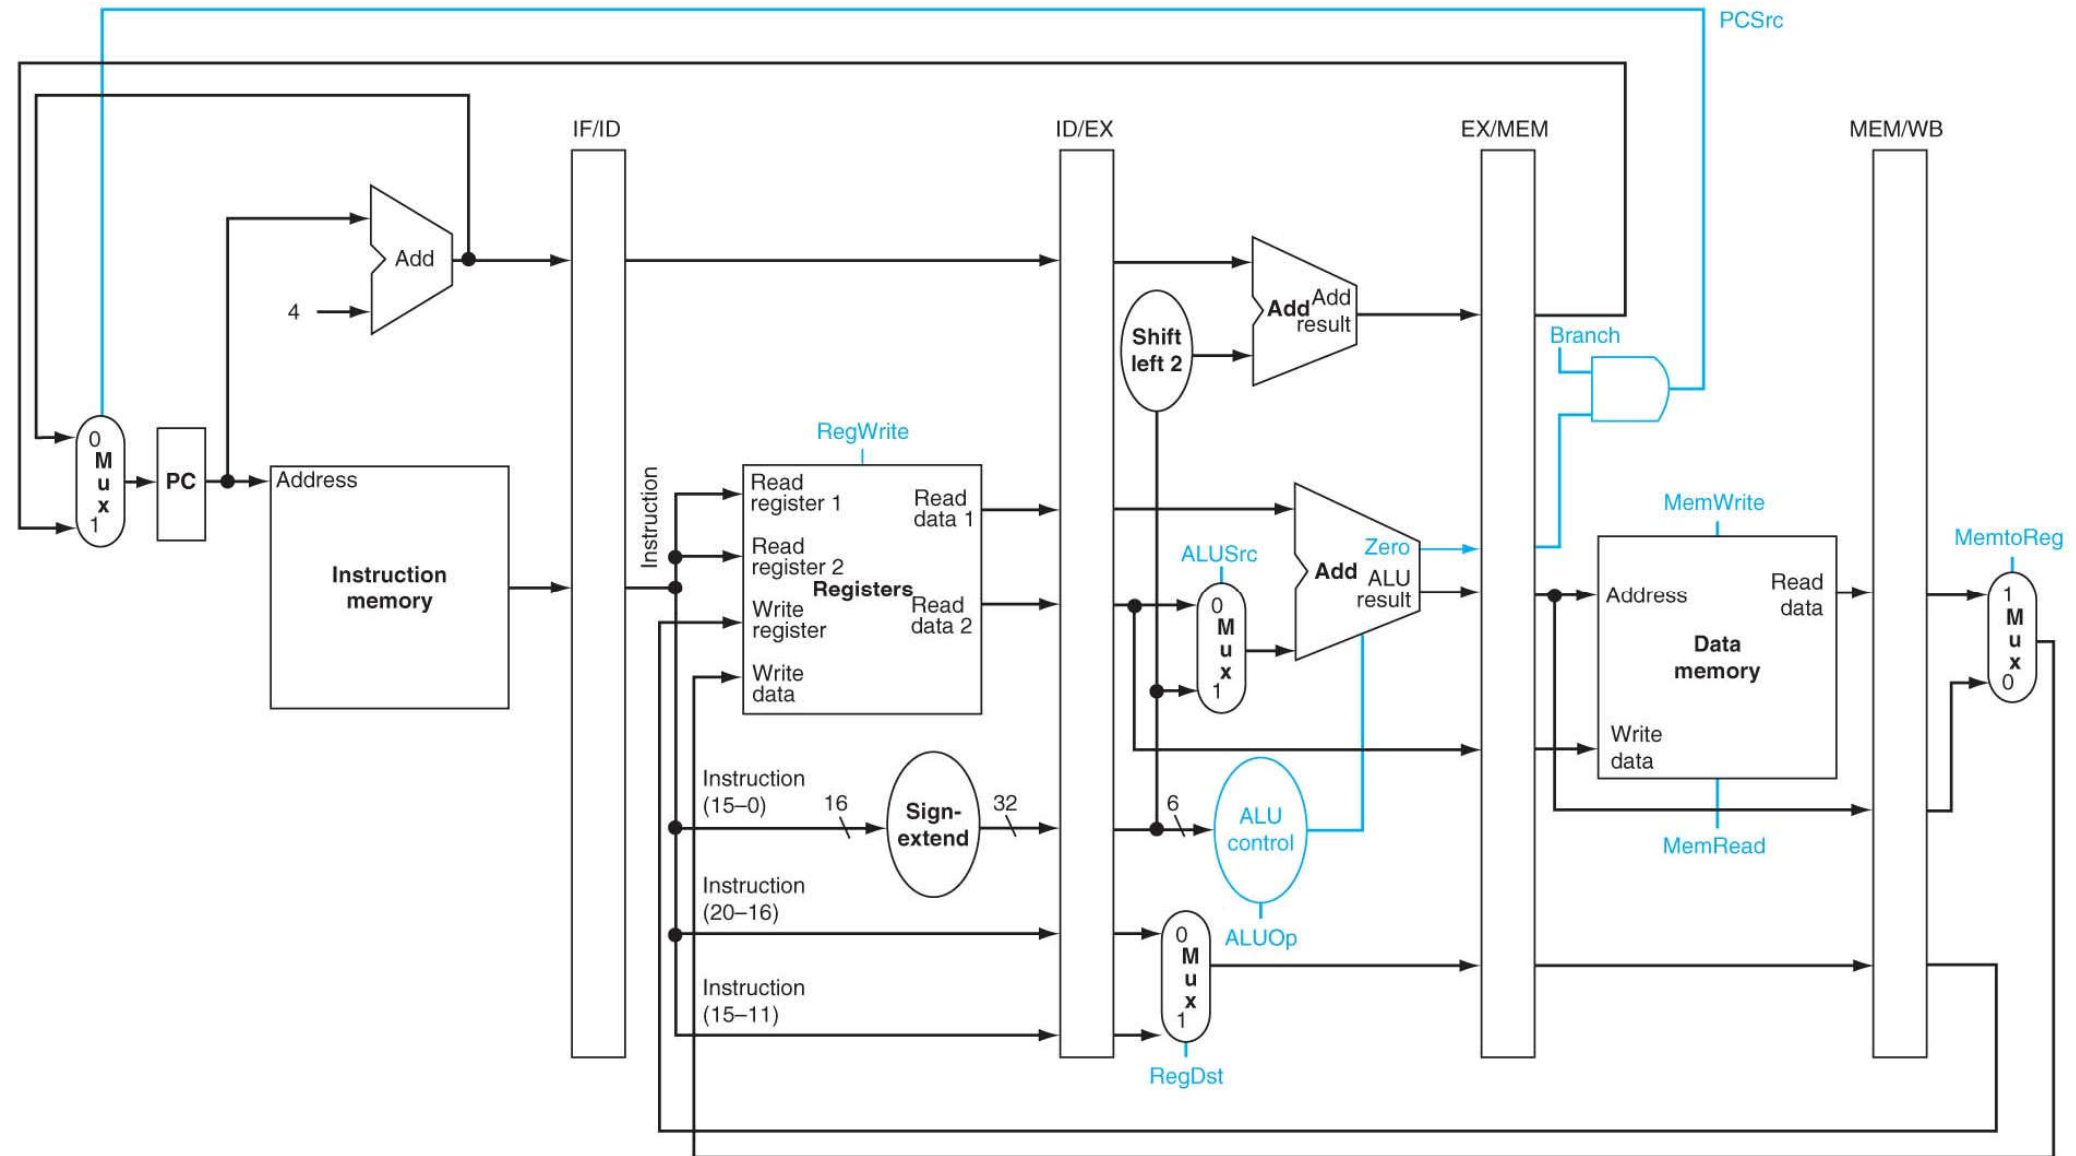
Execute/  
Address Calculation

**sub \$15, \$4, \$1**

**lw \$12, 1000(\$4)**



# The Pipeline, with controls But....






# Pipelined Control


- I told you multicycle control was messy. We would expect pipelined control to be messier.
  - Why?
- But it turns out we can do it with just...
- **Combinational logic!**
  - signals generated **once**, but follow instruction through the pipeline

# Pipelined Control



So, really it is combinational logic and some registers to propagate the signals to the right stage.

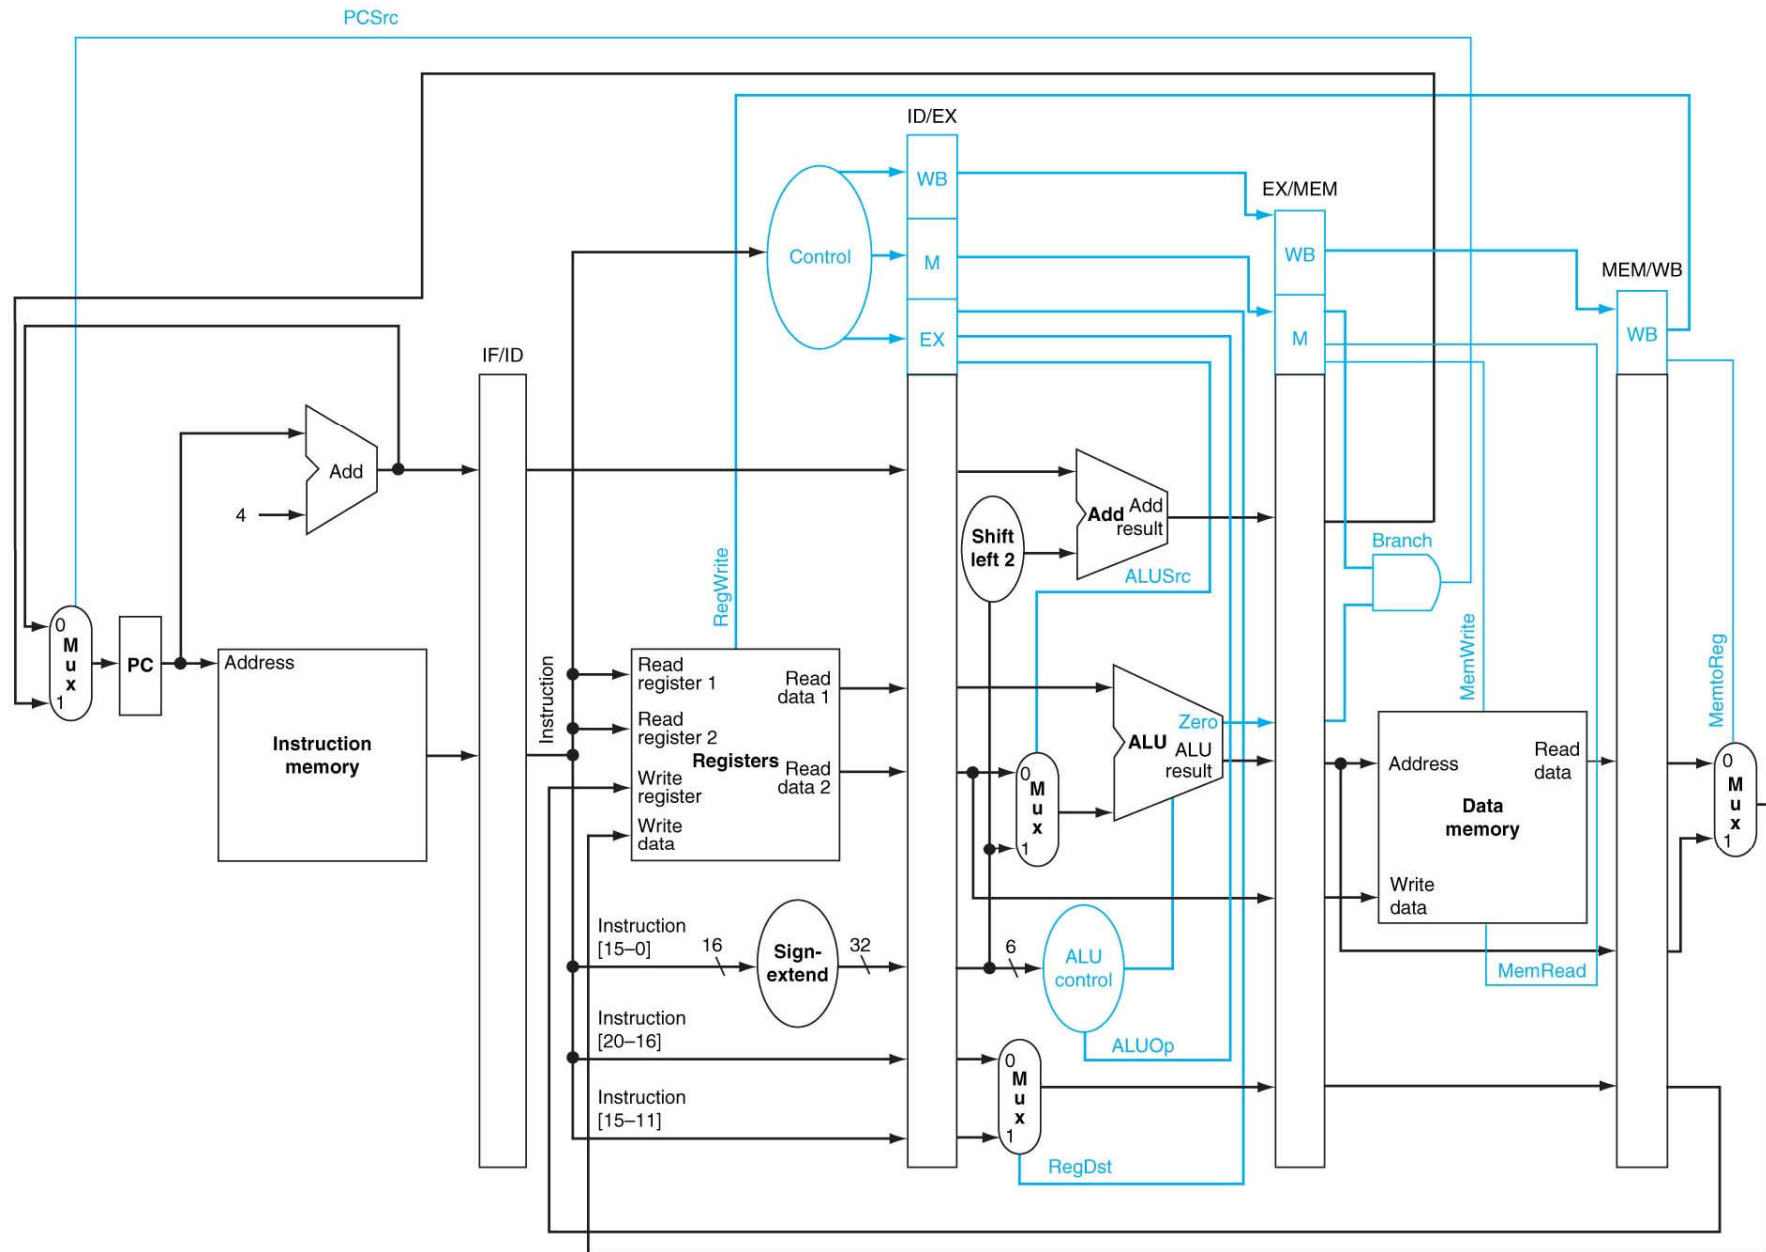
# The Pipeline with Control Logic



# Pipelined Control Signals

	Execution Stage Control Lines				Memory Stage Control Lines			Write Back Stage Control Lines	
Instruction	RegDst	ALUOp1	ALUOp0	ALUSrc	Branch	MemRead	MemWrite	RegWrite	MemtoReg
R-Format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	x	0	0	1	0	0	1	0	x
beq	x	0	1	0	1	0	0	0	x

# The Pipeline with Control Logic



# Is it really that easy?

- What happens when...
  - add \$3, \$10, \$11
  - lw \$8, 1000(\$3)
  - sub \$11, \$8, \$7

# The Pipeline in Execution


**lw \$8, 1000(\$3)**

**add \$3, \$10, \$11**

Execute/  
Address Calculation

Memory Access

Write Back



# The Pipeline in Execution


sub \$11, \$8, \$7

lw \$8, 1000(\$3)

add \$3, \$10, \$11

Memory Access

Write Back





# The Pipeline in Execution

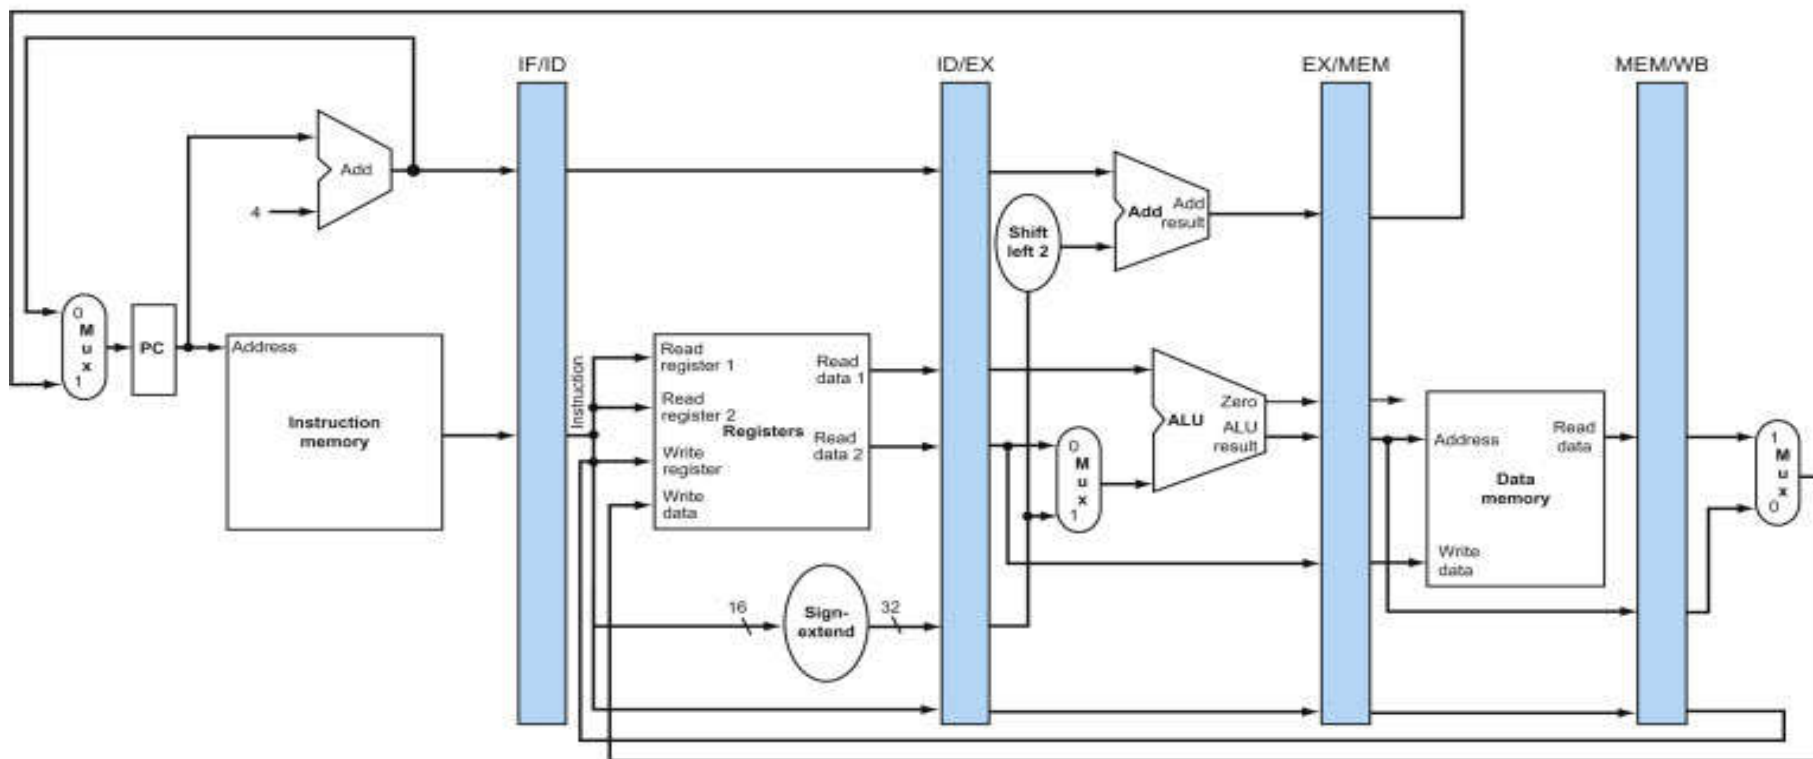
add \$10, \$1, \$2

sub \$11, \$8, \$7

lw \$8, 1000(\$3)

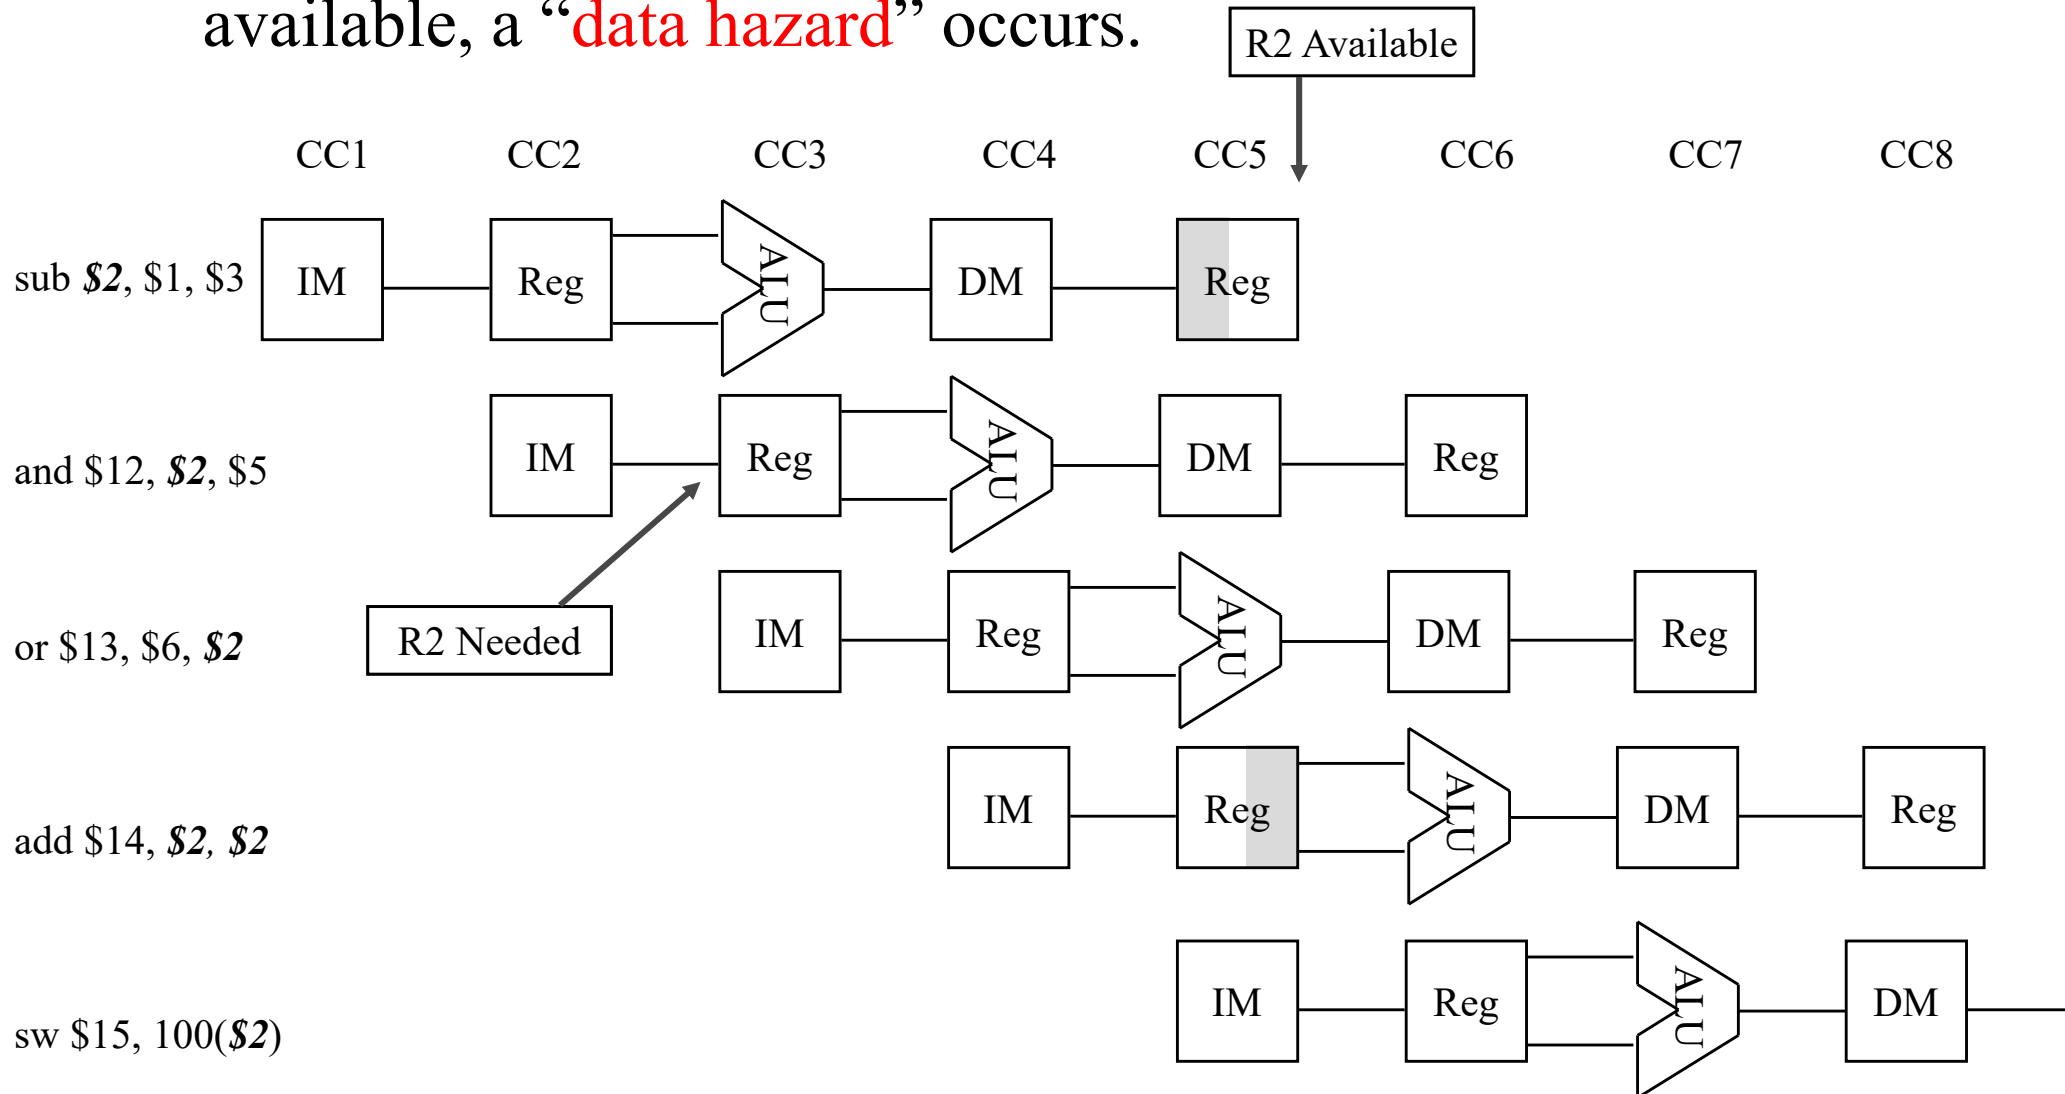
add \$3, \$10, \$11

Write Back



# Data Hazards

- When a result is needed in the pipeline before it is available, a “**data hazard**” occurs.



So... what are we going to do about  
data hazards?

# Pipelining Key Points

- $ET = IC * CPI * CT$
- We achieve high *throughput* without reducing instruction *latency*.
- Pipelining exploits a special kind of parallelism (parallelism between functionality required in different cycles by different instructions).
- Pipelining uses combinational logic to generate (and registers to propagate) control signals.
- Pipelining creates potential hazards.