

Branch Hazards

or

“Which way did she go?”

Control Dependence

- Just as an instruction will be dependent on other instructions to provide its operands (*data dependence*), it will also be dependent on other instructions to determine whether it gets executed or not (*control dependence* or *flow dependence*).
- Control dependences are particularly critical with conditional branches.

add \$5, \$3, \$2

sub \$6, \$5, \$2

beq \$6, \$7, somewhere

and \$9, \$6, \$1

...

somewhere: or \$10, \$5, \$2

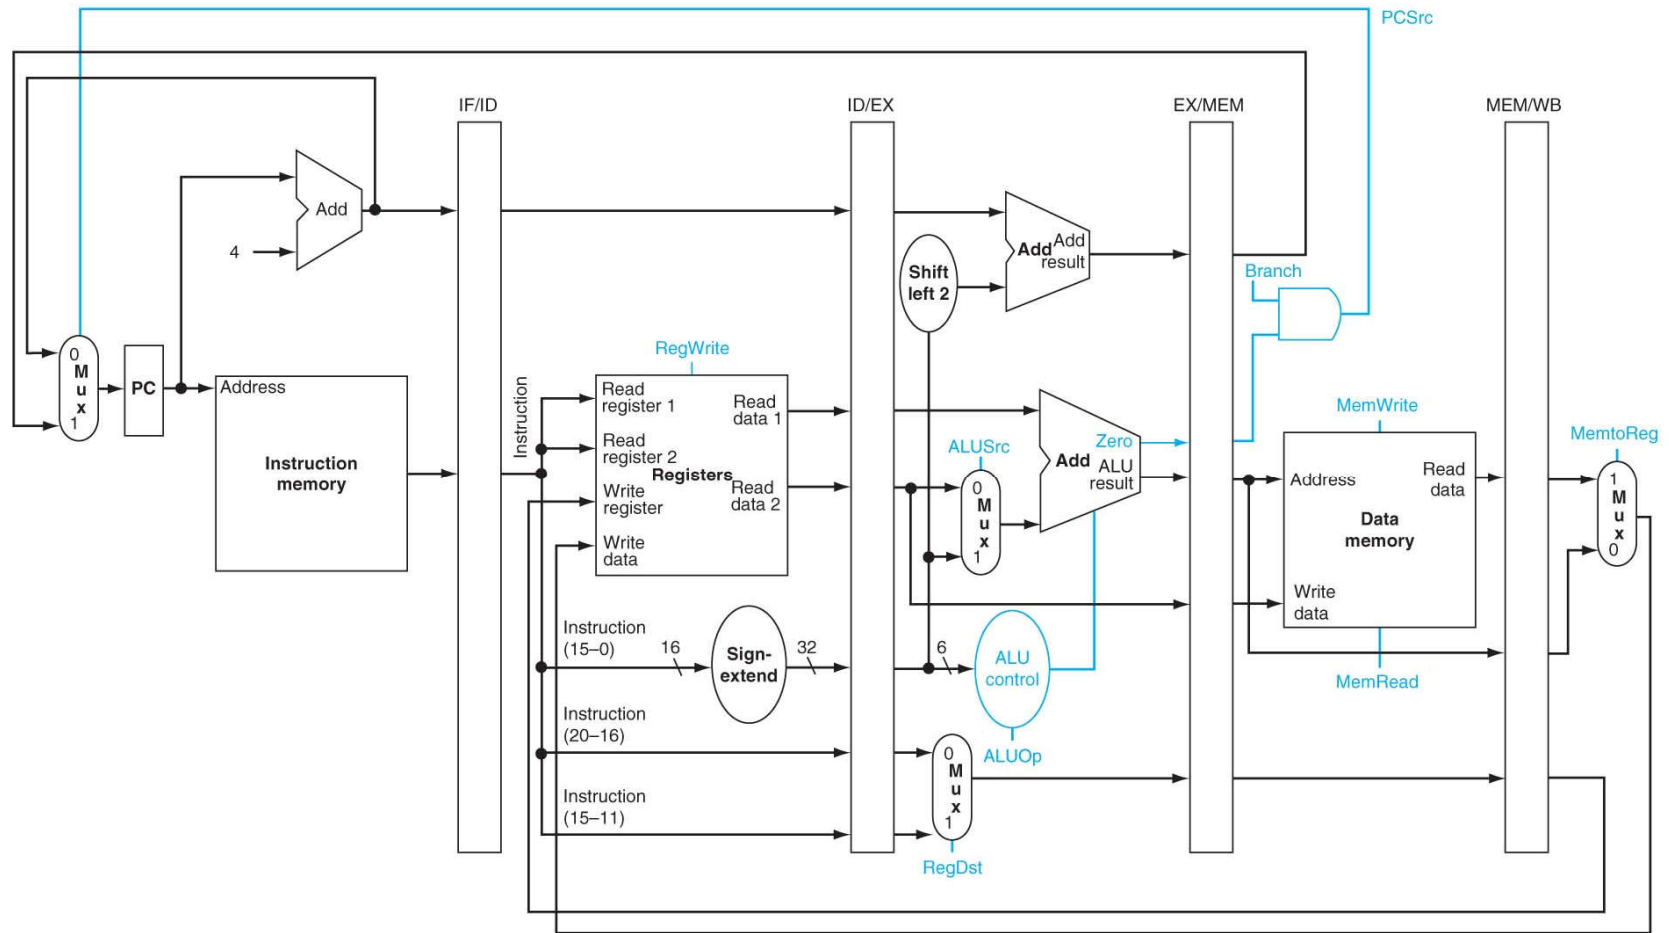
add \$12, \$11, \$9

...

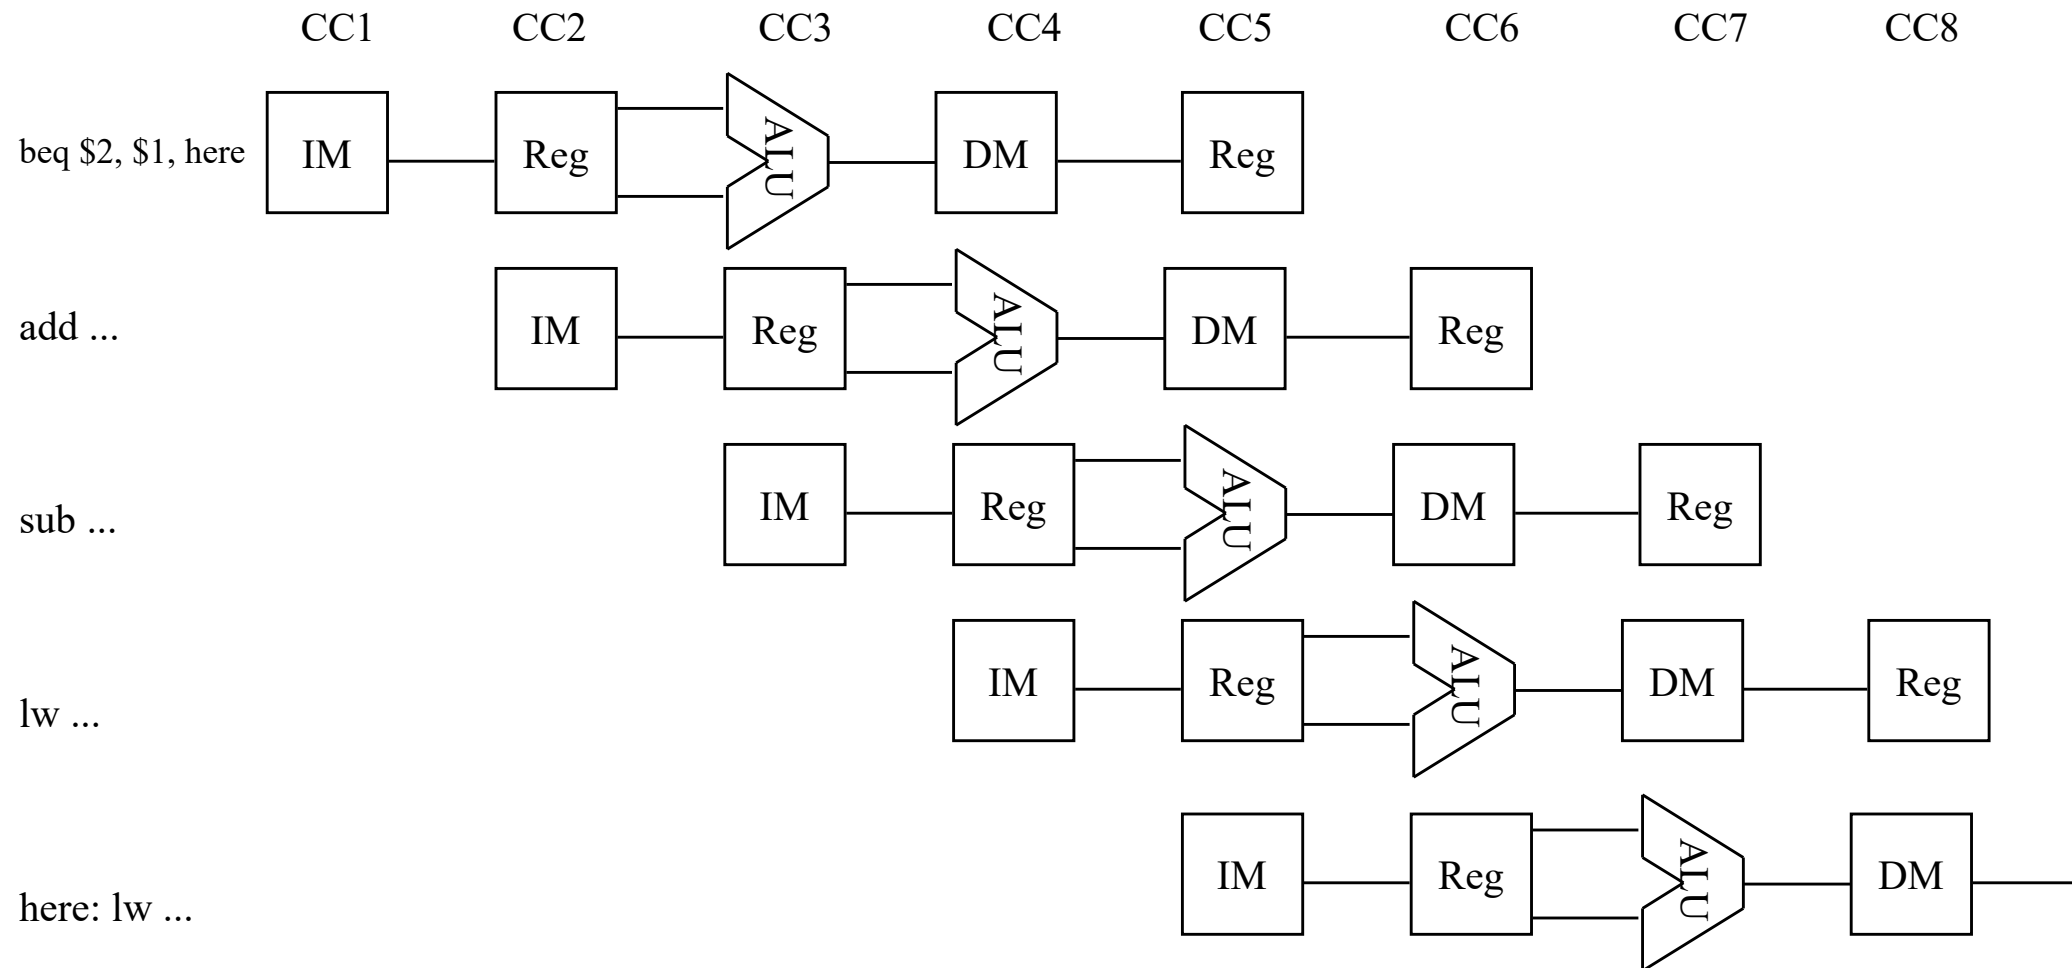
Branch Hazards

- Branch dependences can result in branch hazards (when they are too close to be handled correctly in the pipeline).

Branch Hazards



Branch Hazards



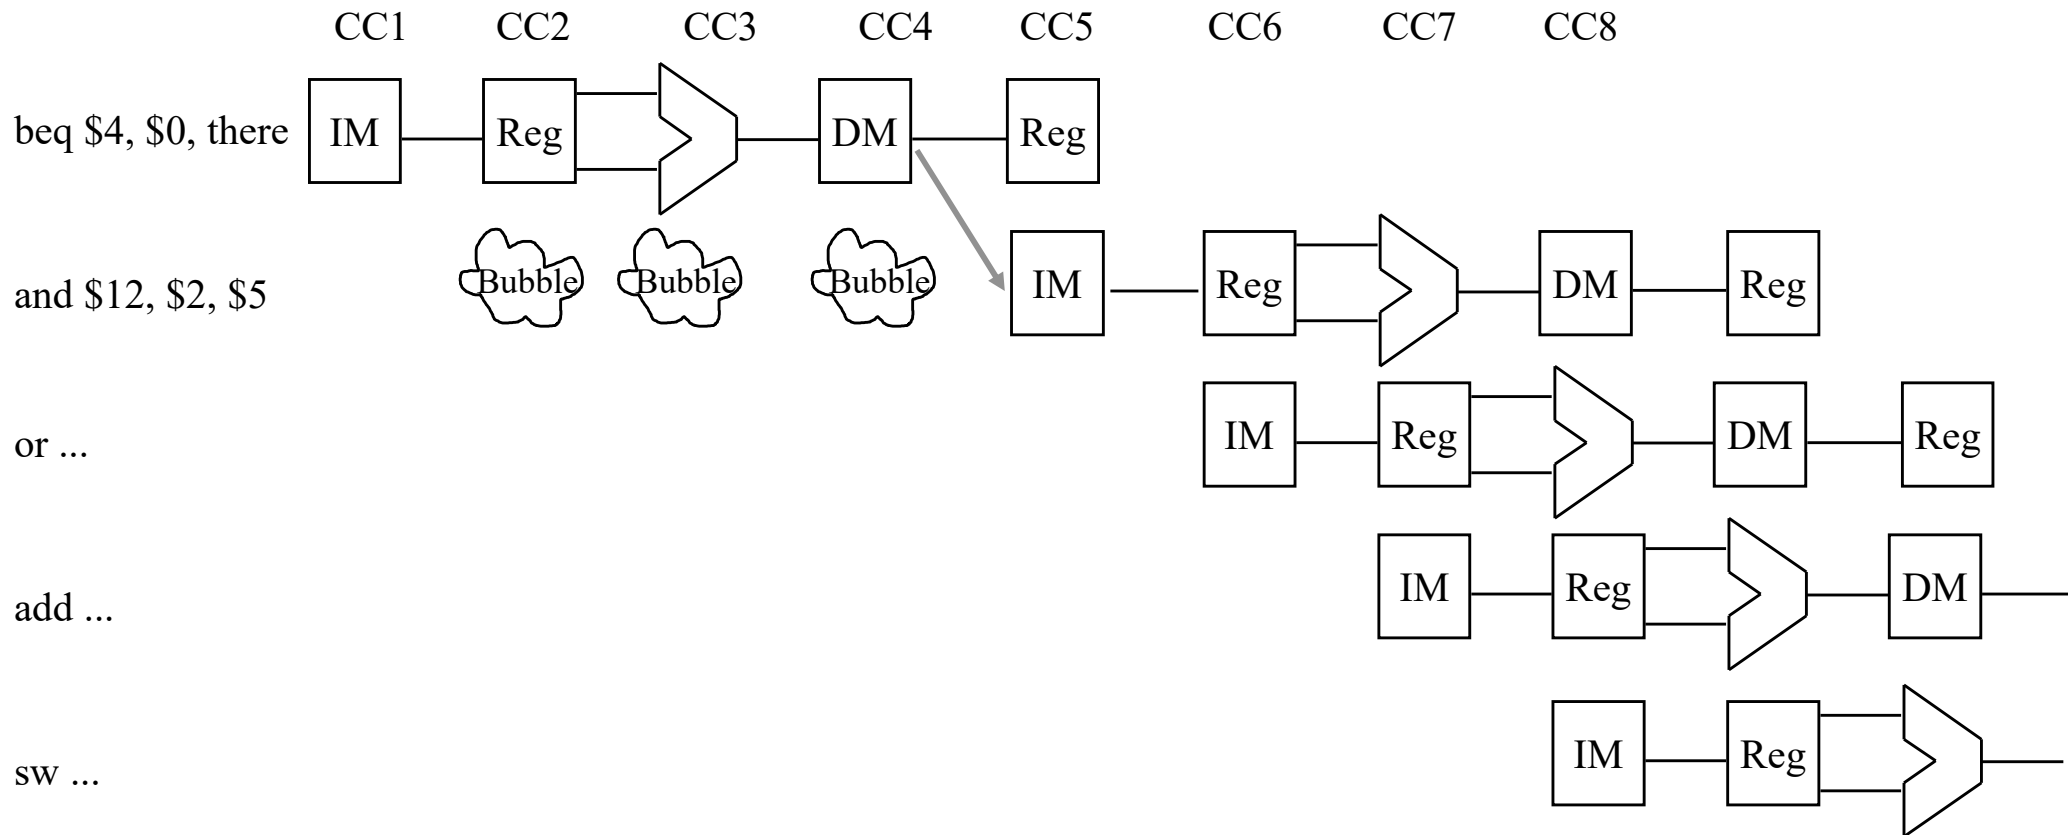
Dealing With Branch Hazards

- Hardware
 - stall until you know which direction
 - reduce hazard through earlier computation of branch direction
 - guess which direction
 - assume not taken (easiest)
 - more educated guess based on history (requires that you know it is a branch before it is even decoded!)

Dealing With Branch Hazards

- Hardware
 - stall until you know which direction
 - reduce hazard through earlier computation of branch direction
 - guess which direction
 - assume not taken (easiest)
 - more educated guess based on history (requires that you know it is a branch before it is even decoded!)
- Hardware/Software/ISA
 - nops, or instructions that get executed either way (delayed branch).

Stalling for Branch Hazards

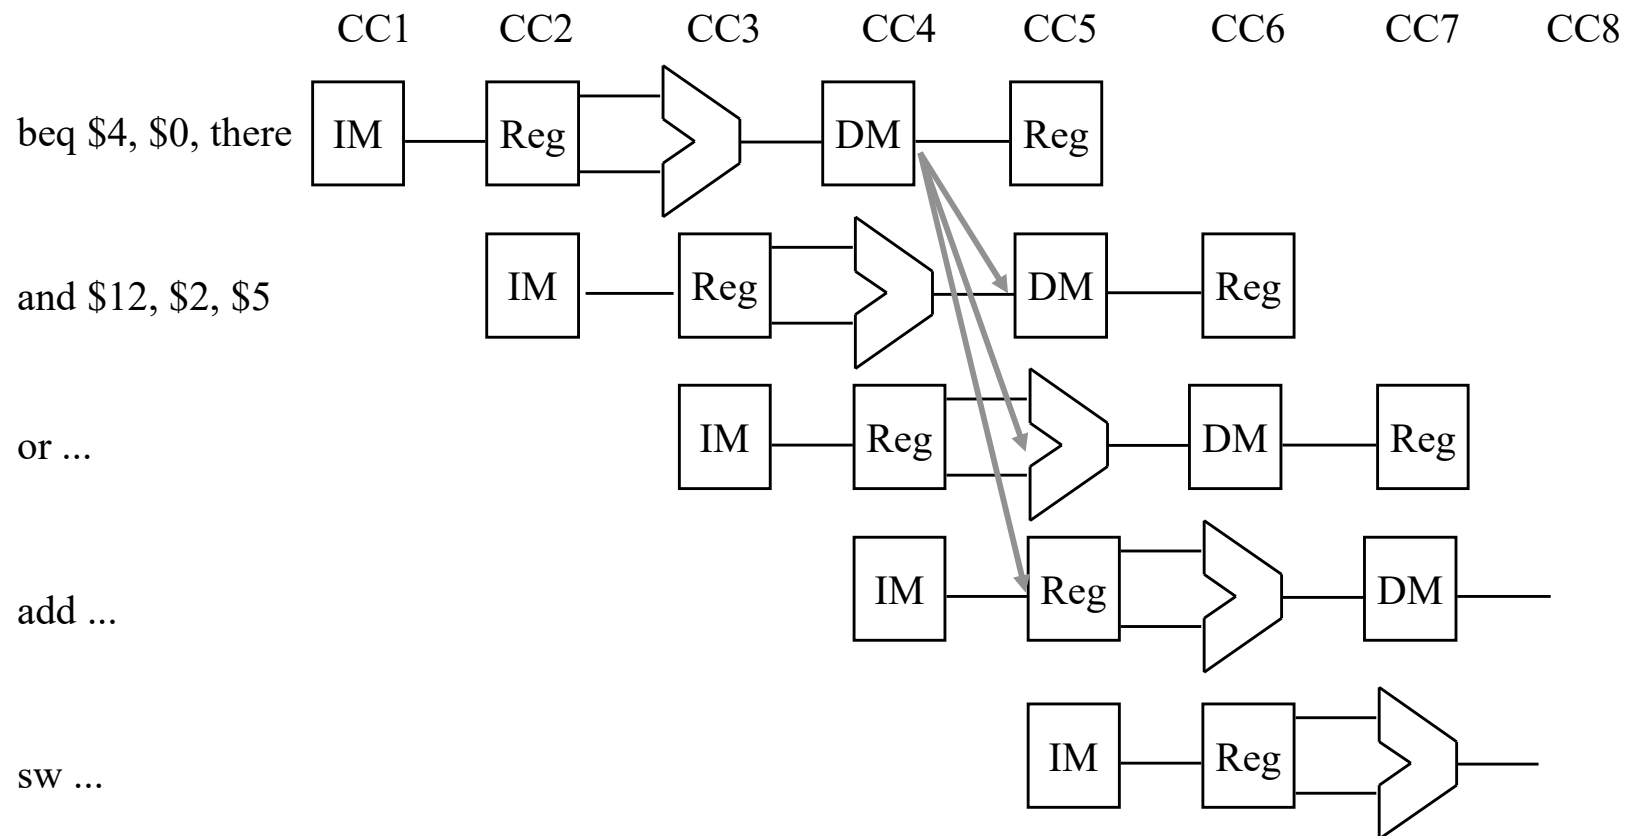


Stalling for Branch Hazards

- Seems wasteful, particularly when the branch isn't taken.
- Makes all branches cost 4 cycles.
- Also, requires you know that it's a branch before you decode it (several of our solutions have this problem)

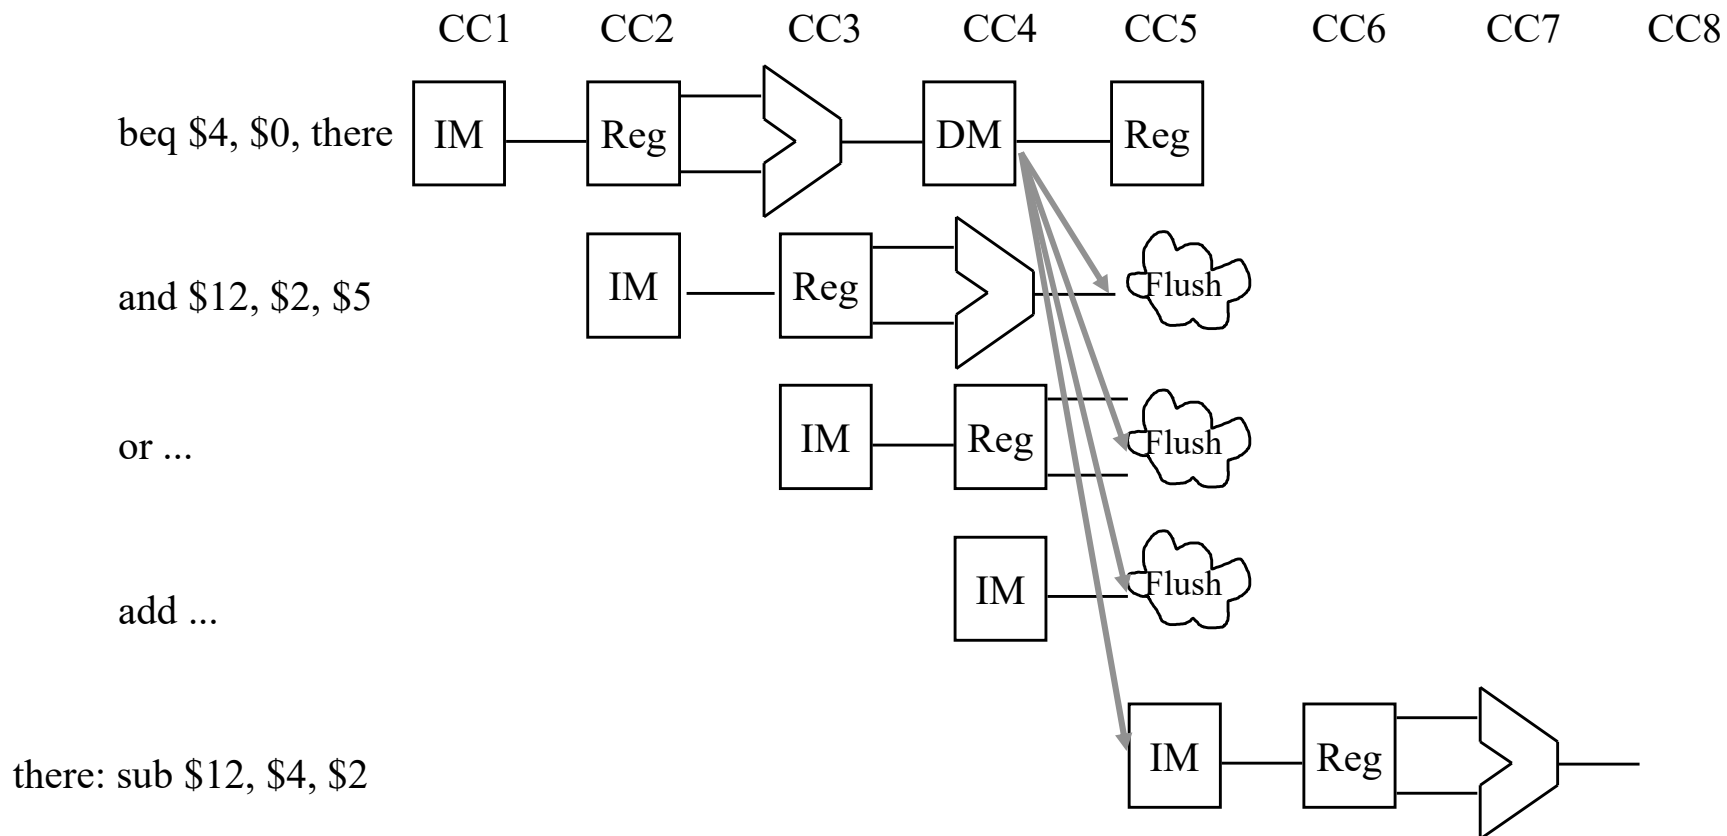
Assume Branch *Not Taken*

- works pretty well when you're right



Assume Branch *Not Taken*

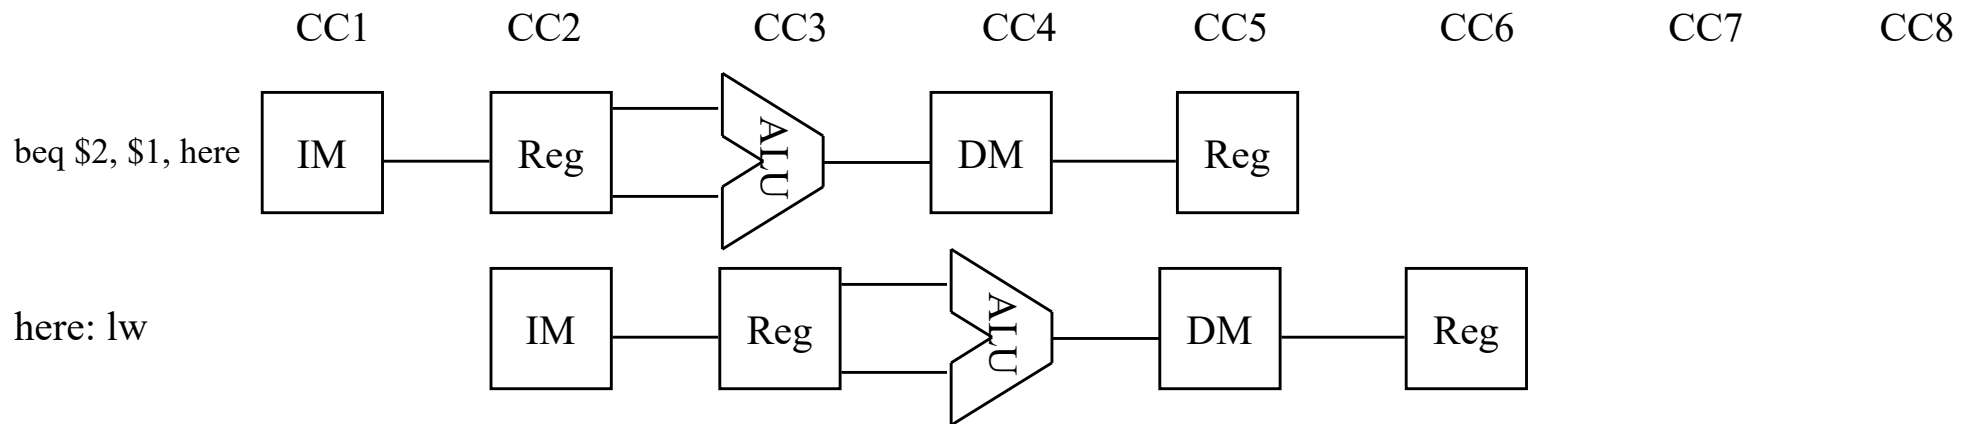
- same performance as stalling when you're wrong



Assume Branch *Not Taken*

- Performance depends on percentage of time you guess right.
- Flushing an instruction means to prevent it from changing any permanent state (registers, memory, PC).
 - sounds a lot like a bubble...
 - But notice that we need to be able to insert those bubbles later in the pipeline

Branch Hazards – Predicting Taken?

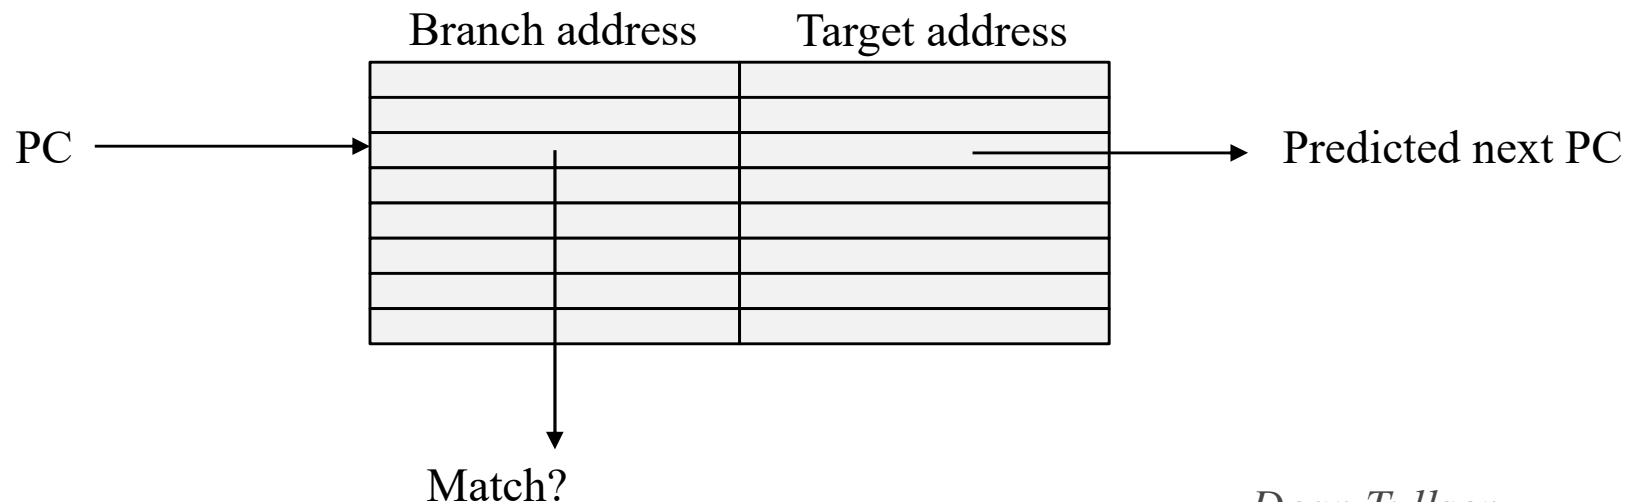


Required information to predict Taken:

1. Whether an instruction is a branch (before decode)
2. The target of the branch
3. The outcome of the branch condition

Branch Target Buffer

- Keeps track of the PCs of recently seen branches and their targets.
- Consult during Fetch (in parallel with Instruction Memory read) to determine:
 - Is this a branch?
 - If so, what is the target

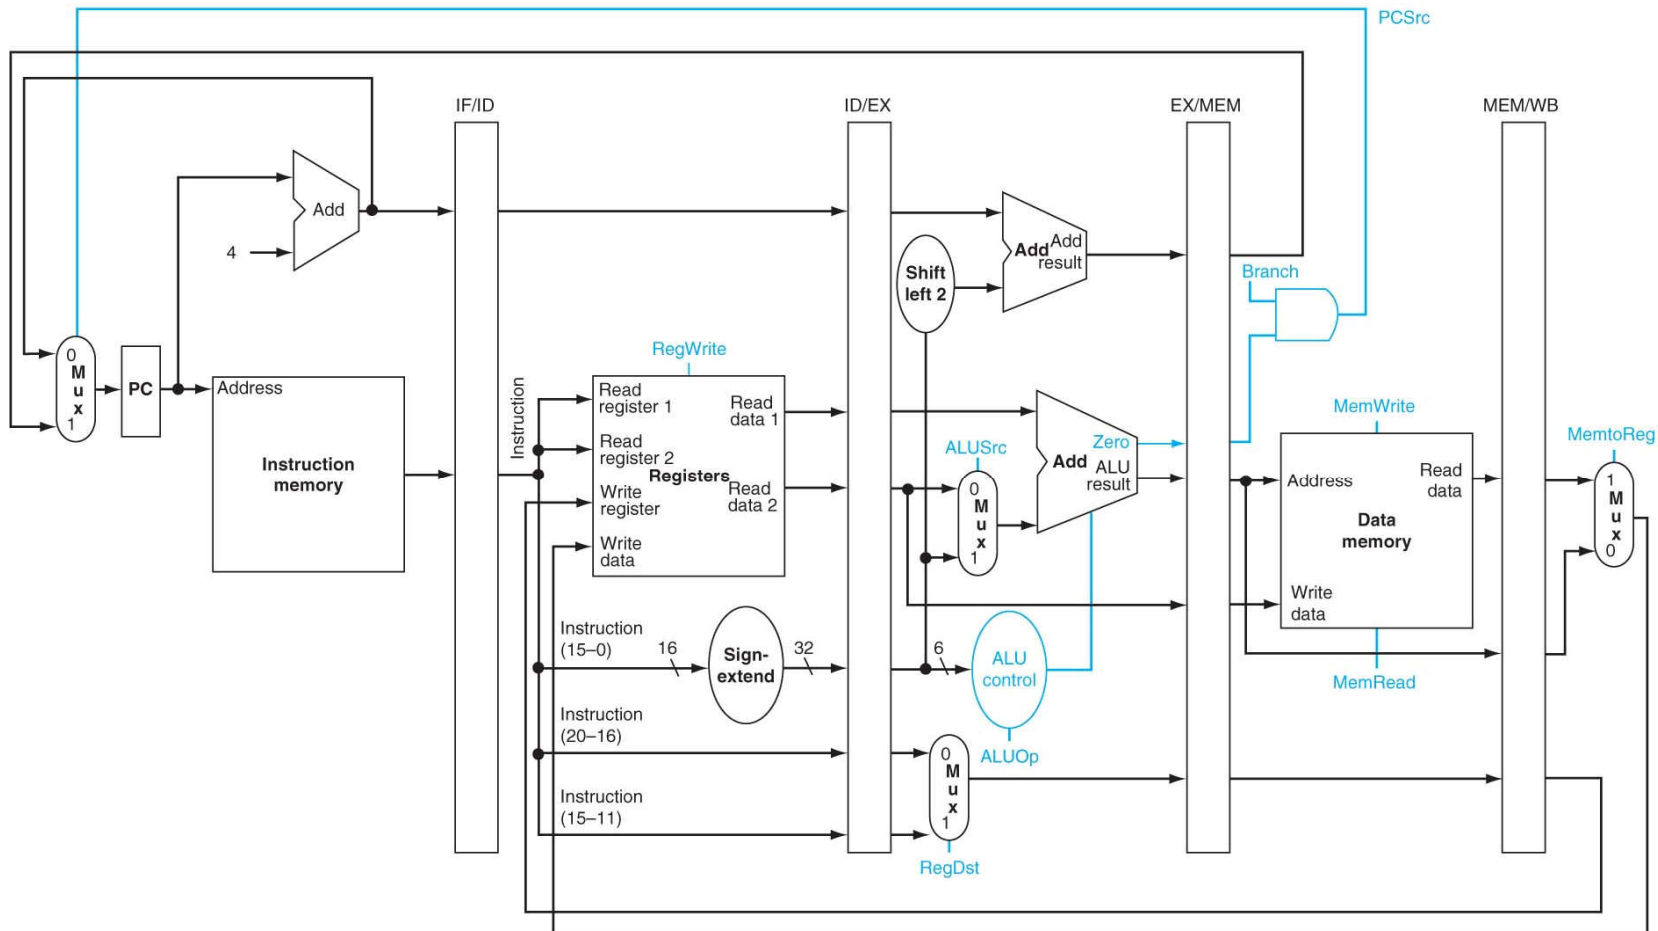


So, just to set the stage...

MIPS R2000

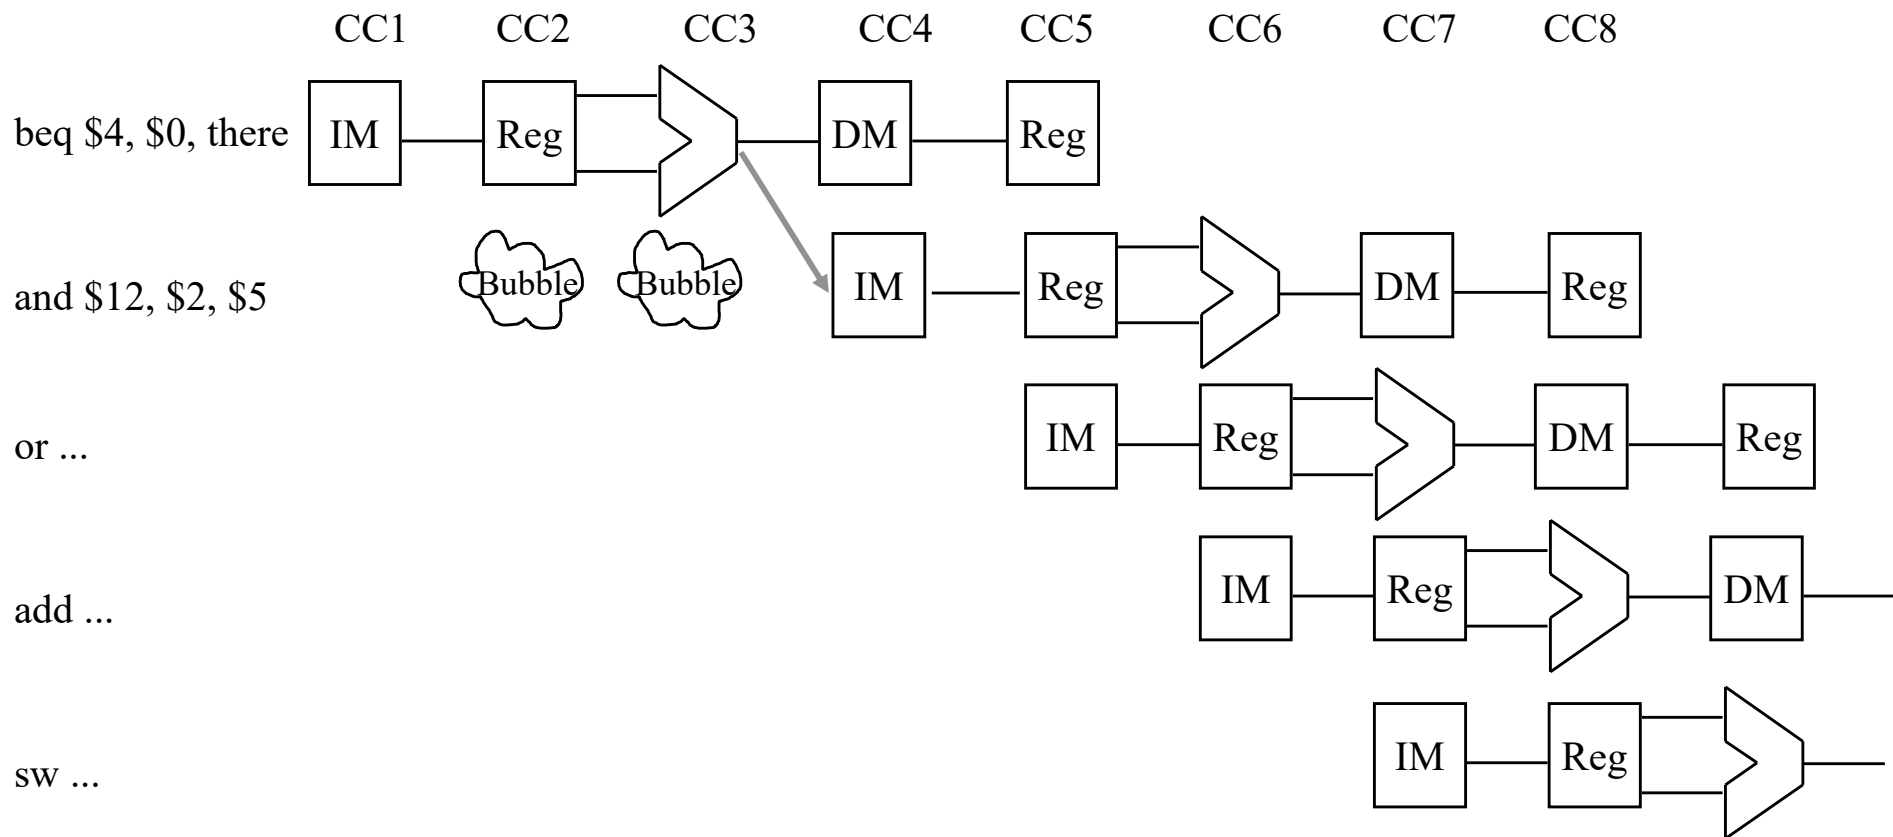
- We are going to walk through the solutions they used for the initial commercial MIPS pipeline.
- Just be aware, many of those do not transfer to modern architectures.
- But this will lay the foundation for modern techniques, which we'll look at in some detail.

Reducing the Branch Delay

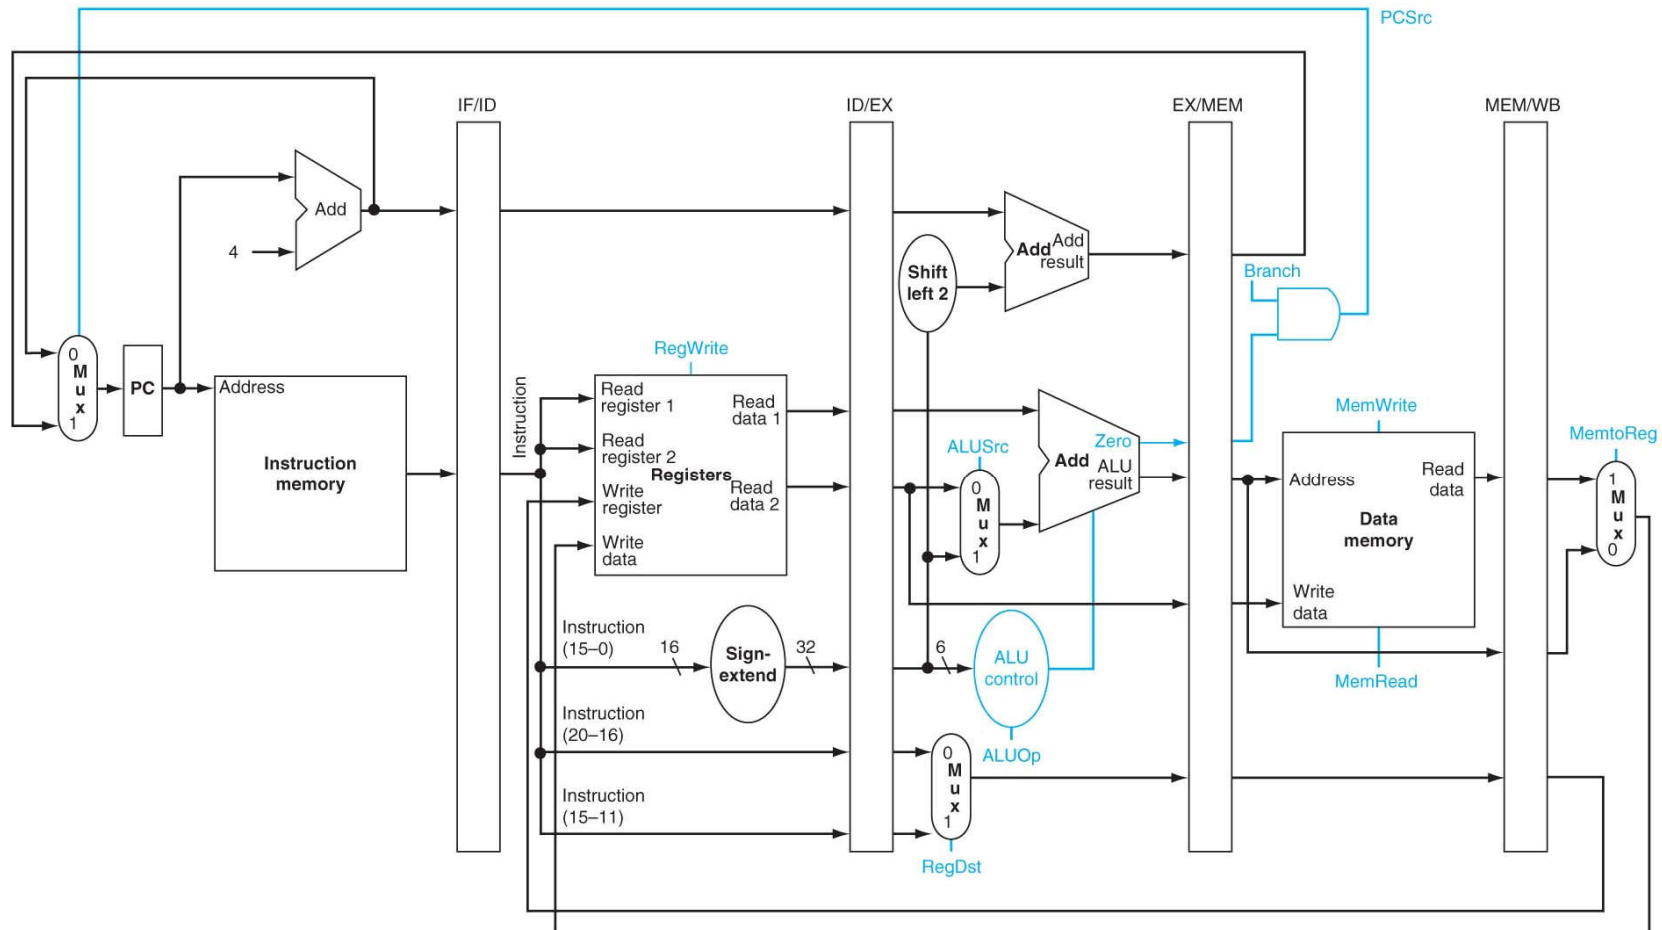


- can easily get to 2-cycle stall

Stalling for Branch Hazards

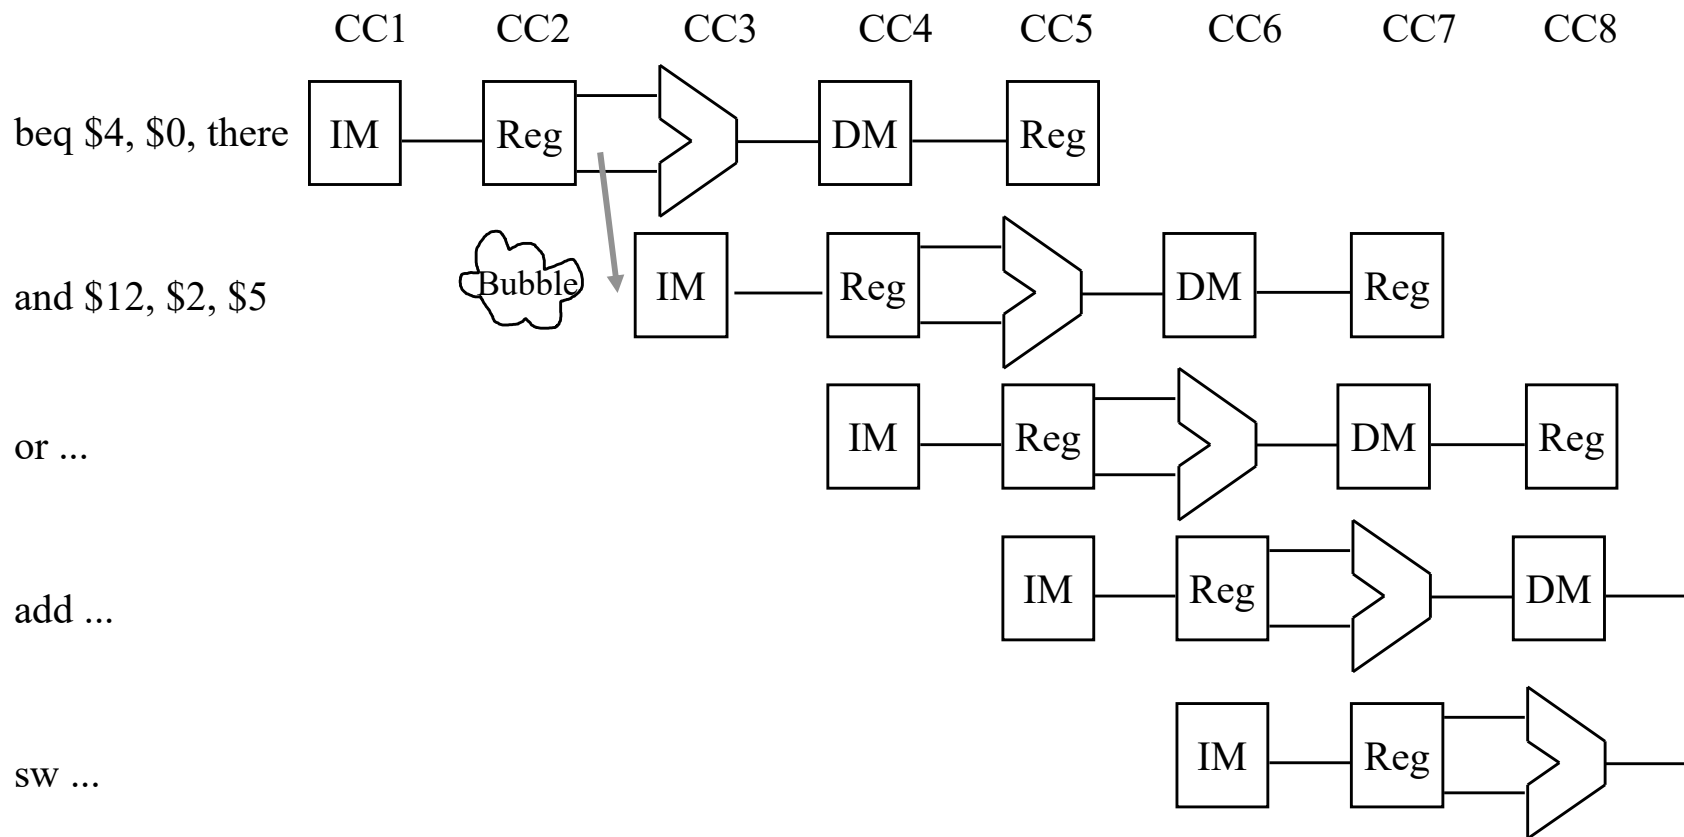


Reducing the Branch Delay

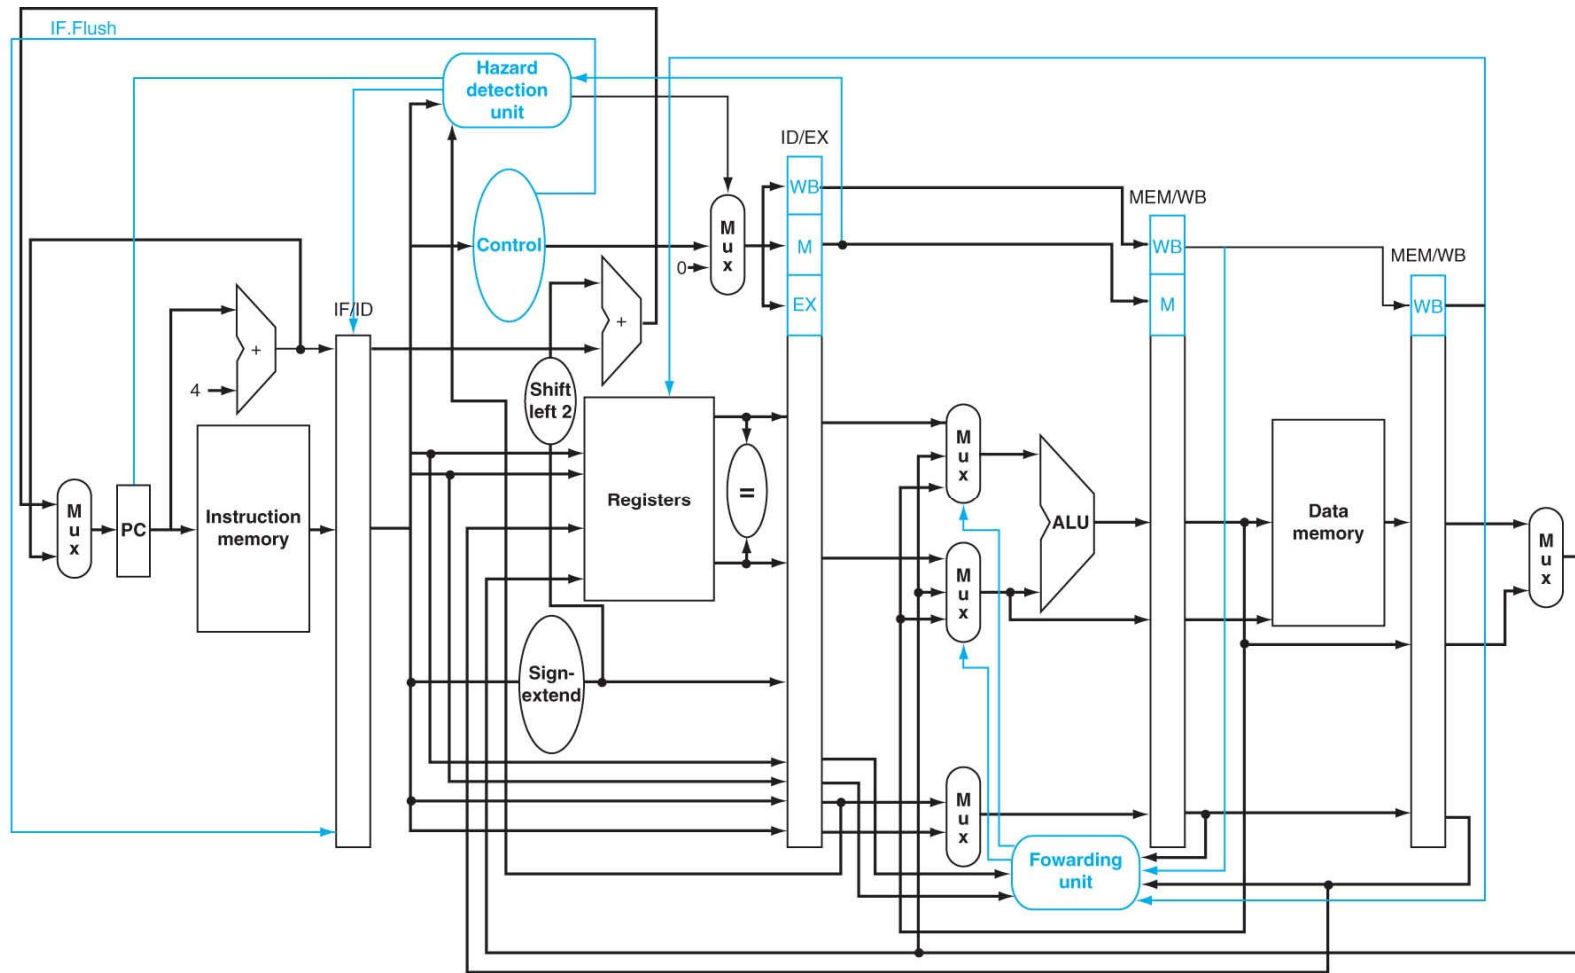


- Harder, but possible, to get to 1-cycle stall

Stalling for Branch Hazards



The Pipeline with flushing for taken branches

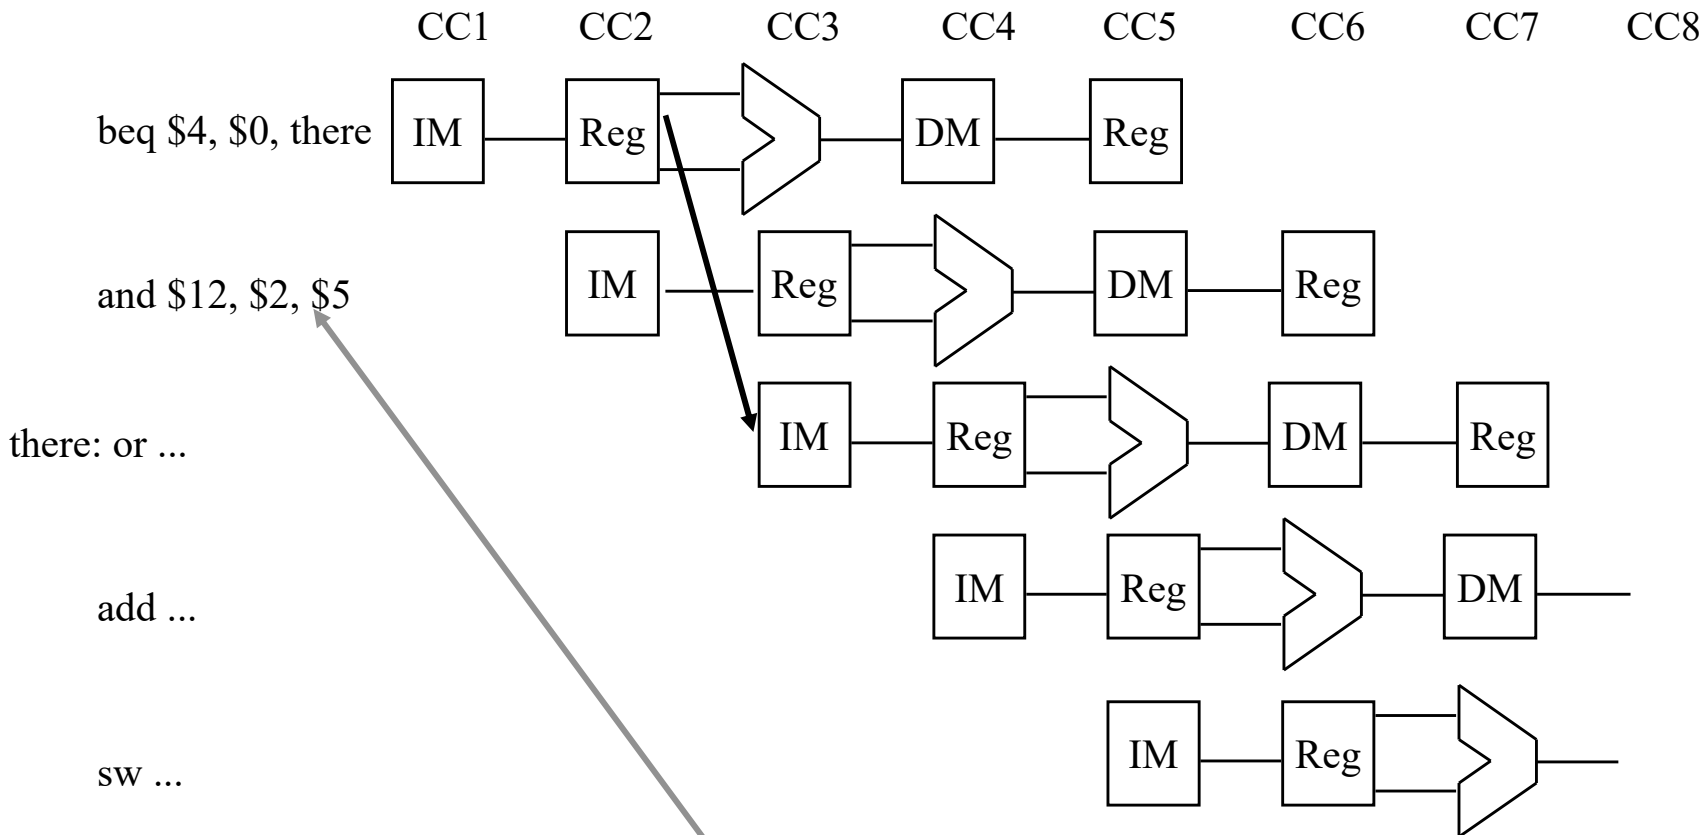


- Notice the IF/ID flush line added.
- Not our final design yet, so don't need to study this too hard.
- MIPS ISA choice -> beq, blt, ...

Eliminating the Branch Stall

- There's no rule that says we have to see the effect of the branch immediately. Why not wait an extra instruction before branching?
- The original SPARC and MIPS processors each used a single *branch delay slot* to eliminate single-cycle stalls after branches.
- The instruction after a conditional branch is *always executed* in those machines, regardless of whether the branch is taken or not!

Branch Delay Slot



Branch delay slot instruction (next instruction after a branch) is executed even if the branch is taken.

Filling the branch delay slot

- The branch delay slot is only useful if you can find something to put there.
- If you can't find anything, you must put a *nop* to insure correctness.
- Where do we find instructions to fill the branch delay slot?
 -
 -
 -

Filling the branch delay slot

```
1 add $5, $3, $7
2 add $9, $1, $3
3 sub $6, $1, $4
4 and $7, $8, $2
5 beq $6, $7, there
  nop /* branch delay slot */
6 add $9, $1, $4
7 sub $2, $9, $5
  ...
there:
8 mult $2, $10, $11
  ...
```

- Which instructions could be used to replace the nop?

Branch Delay Slots

- This works great for this implementation of the architecture, but becomes a permanent part of the ISA.
- What about the MIPS R10000, which has a 5-cycle branch penalty, and executes 4 instructions per cycle??
- What about the Pentium 4, which has a 21-cycle branch penalty and executes up to 3 instructions per cycle???

Early resolution of branch + branch delay slot

- Worked well for MIPS R2000 (the 5-stage pipeline MIPS)
- Early resolution doesn't scale well to modern architectures
 - Better to always have execute happen in execute
 - Forwarding into branch instruction?
- Branch delay slot
 - Doesn't solve the problem in modern pipelines
 - Still in ISA, so have to make it work even though it doesn't provide any significant advantage.
 - Violates important general principal – (unless you really only want a single generation of your product) do not expose current technology limitations to the ISA!!

Okay, then...

- What do we do in modern architectures???

Branch Prediction

- Always assuming the branch is not taken is a crude form of *branch prediction*.

Branch Prediction

- Always assuming the branch is not taken is a crude form of *branch prediction*.
- What about loops that are *taken* 95% of the time?
 - we would like the option of assuming *not taken* for some branches, and *taken* for others, depending on ???


Branch Prediction

- Historically, two broad classes of branch predictors:
- *Static predictors* – for branch B, always make the same prediction.
- *Dynamic predictors* – for branch B, make a new prediction every time the branch is fetched.
- Tradeoffs?
- Modern CPUs all have sophisticated *dynamic* branch prediction.

Dynamic Branch Prediction

- What information is available to make an intelligent prediction?

Branch Prediction



```
for (i=0;i<10;i++) {
```

```
...
```

```
...
```

```
}
```



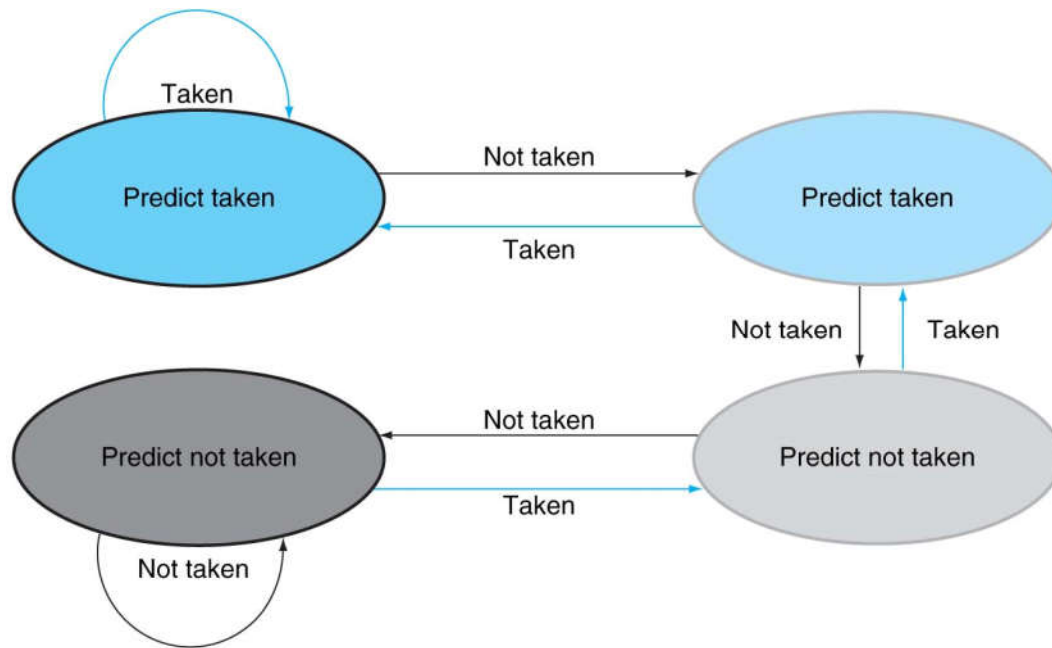
```
...
```

```
...
```

```
add $i, $i, #1
```

```
beq $i, #10, loop
```

Two-bit predictors give better loop prediction



```
for (i=0;i<10;i++) {  
...  
...  
}
```




```
...  
...  
add $i, $i, #1  
beq $i, #10, loop
```

This state machine also referred to as a *saturating counter* – it counts down (on *not takens*) to 00 or up (on *takens*) to 11, but does not wrap around.

Branch History Table (*bimodal* predictor)

- has limited size
- 2 bits by N (e.g. 4K)
- uses low bits of branch address to choose entry




- what about even/odd branch?

bimodal predictor

- For the following loop, what will be the prediction accuracy of the bimodal predictor for the conditional branch that closes the loop?


```
for (i=0; i< 2; i++) //two iterations per loop
{ z = ... }
```

(let's make this a freeform answer poll question)



Selection	Accuracy
A	100%
B	50%
C	0%
D	Maybe 0%, maybe 50%
E	other


2-bit bimodal prediction accuracy



Is this good enough?

Can We Do Better?

- Can we get more information dynamically than just the recent bias of this branch?
- We can look at patterns (2-level *local* predictor) for a particular branch.
 - last eight branches 00100100, then it is a good guess that the next one is “1” (taken)



- even/odd branch?

Can We Do Better?

- *Correlating Branch Predictors* also look at other branches for clues

if (i == 0)

...


if (i > 7)

...

- Typically use two indexes
 - Global history register --> history of last m branches (e.g., 0100011)
 - branch address

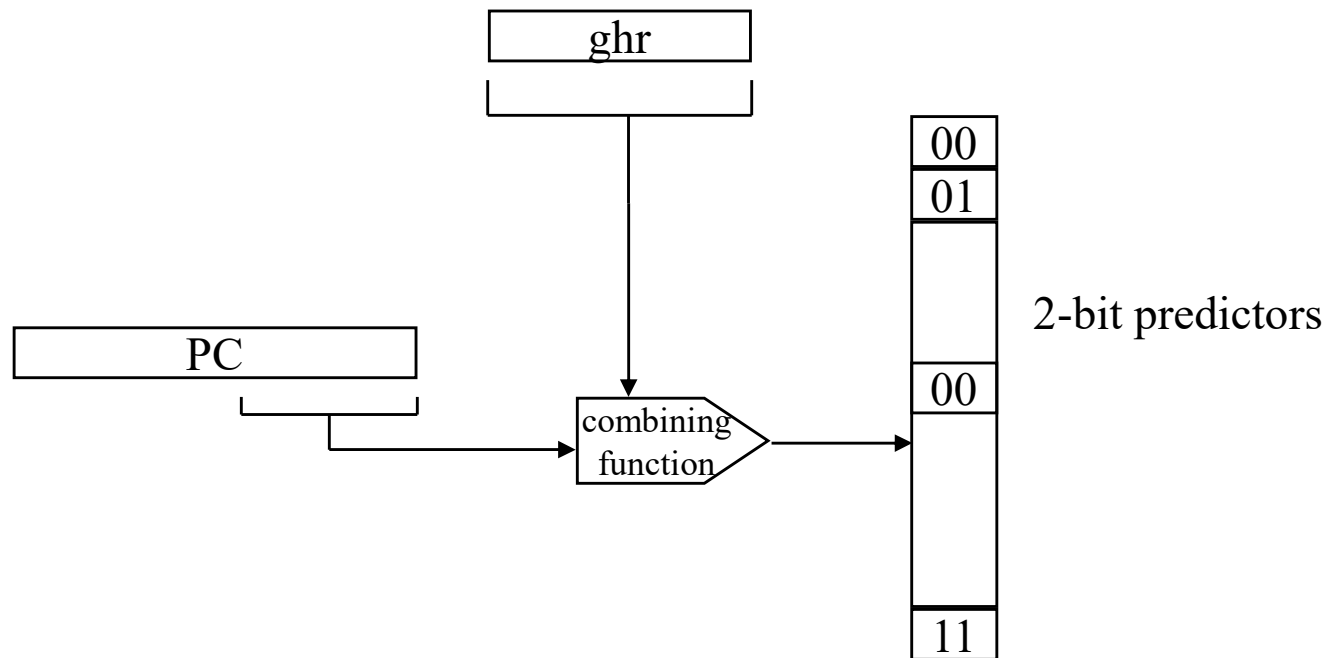
Correlating Branch Predictors

- The *global history register (ghr)* is a shift register that records the last n branches (of any address) encountered by the processor.



Two-level correlating branch predictors

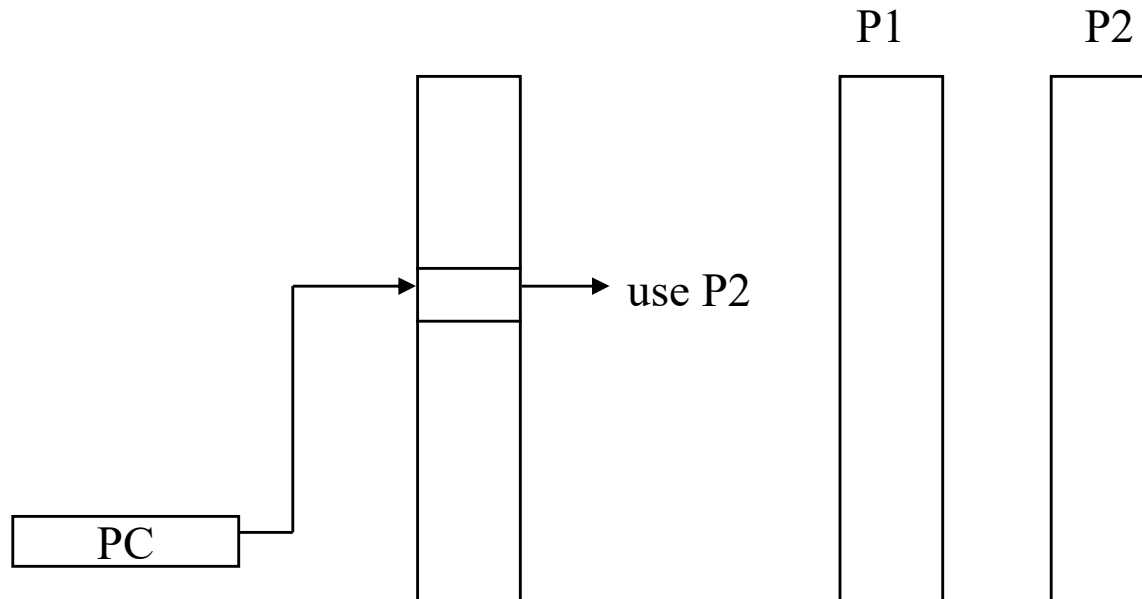
- Can use both the PC address and the GHR




- Most common – *gshare*: use xor as the combining function.

Are we happy yet????

- *Combining branch predictors* use multiple schemes and a voter to decide which one typically does better for *that* branch.



Compaq/Digital Alpha 21264



Aliasing in Branch Predictors

- Branch predictors will always be of finite size, while code size is relatively unlimited.


Aliasing in Branch Predictors

- Branch predictors will always be of finite size, while code size is relatively unlimited.
- What happens when (in the common case) there are more branches than entries in the branch predictor?


Aliasing in Branch Predictors

- Branch predictors will always be of finite size, while code size is relatively unlimited.
- What happens when (in the common case) there are more branches than entries in the branch predictor?
- We call these conflicts *aliasing*.
- We can have negative aliasing (when biases are different) or neutral aliasing (biases same). Positive aliasing is unlikely.

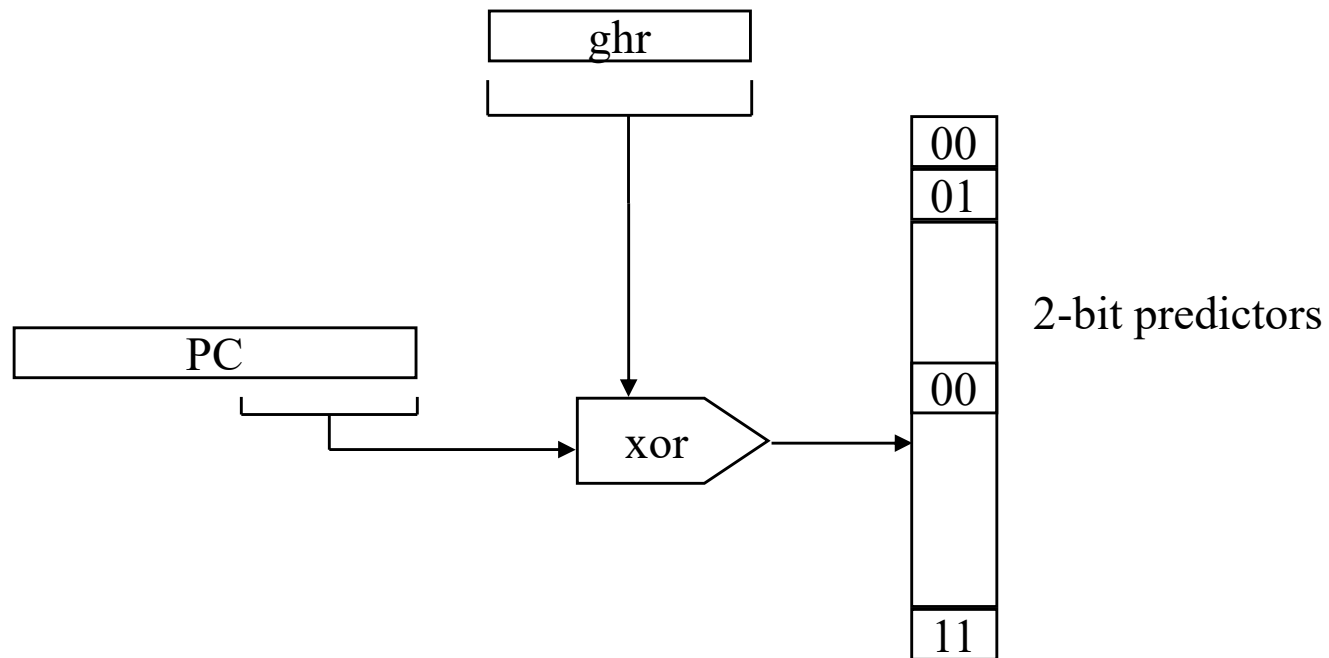
Bimodal aliasing



Local Predictor Aliasing



Gshare aliasing



Branch Prediction

- Latest branch predictors significantly more sophisticated, using more advanced correlating techniques, larger structures, and in some cases using AI techniques.
- Remember from earlier....
 - Presupposes what two pieces of information are available at *fetch time*?
 -
 -
 - Branch Target Buffer supplies this information.

MIPS R2000

- Okay, at this point, we've essentially recreated the MIPS R2000 (ignoring the branch prediction slides – that all came later).
- R2000
 - 5 stage pipeline
 - Forwarding wherever possible
 - 1 cycle load-use hazard
 - Branches resolved in ID stage
 - Branch Delay Slot incorporated in the ISA.
- So what kind of performance should we expect?

Pipeline performance

```
loop: lw $15, 1000($2)
      add $16, $15, $12
      lw $18, 1004($2)
      add $19, $18, $12
      beq $19, $0, loop:
      nop
```

What is the steady-state CPI of this code? Assume branch taken many times. Assume 5-stage pipeline, forwarding, early branch resolution, BDS. (always assume this architecture if not given the details)

Can we improve this?

Putting it all together.

For a given program on our 5-stage MIPS pipeline processor:
20% of insts are loads, 50% of instructions following a load
are arithmetic instructions depending on the load
20% of instructions are branches. We manage to fill 80% of
the branch delay slots with useful instructions.

What is the CPI of your program?

Given our 5-stage MIPS pipeline – what is the steady state CPI for the following code? Assume the branch is taken thousands of times.

```
Loop: lw r1, 0(r2)
      add r2, r3, r4
      sub r5, r1, r2
      beq r5, $zero, Loop
      nop
```

IF = 200ps
ID = 100ps
EX = 150ps
M = 200ps
WB = 100ps

Hardware engineers determine these to be the execution times per stage of the MIPS 5-stage pipeline processor. Consider splitting IF and M into 2 stages each. (So IF1 IF2 and M1 M2.) For this code (assume branch usually taken):

```
Loop: lw r1, 0(r2)
      add r2, r3, r4
      sub r5, r1, r2
      beq r5, $zero, Loop
      nop
```

What would be the impact of the new 7-stage pipeline compared to the original 5-stage MIPS pipeline.. Assume the pipeline has forwarding where necessary, predicts branch not taken, and resolves branches in ID.

7-stage Pipeline

```
Loop: lw r1, 0(r2)
      add r2, r3, r4
      sub r5, r1, r2
      beq r5, $zero, Loop
      nop
```


Control Hazards -- Key Points

- Control (or branch) hazards arise because we must fetch the next instruction before we know if we are branching or where we are branching.
- Control hazards are detected in hardware.
- We can reduce the impact of control hazards through:
 - early detection of branch address and condition
 - *branch prediction*
 - branch delay slots